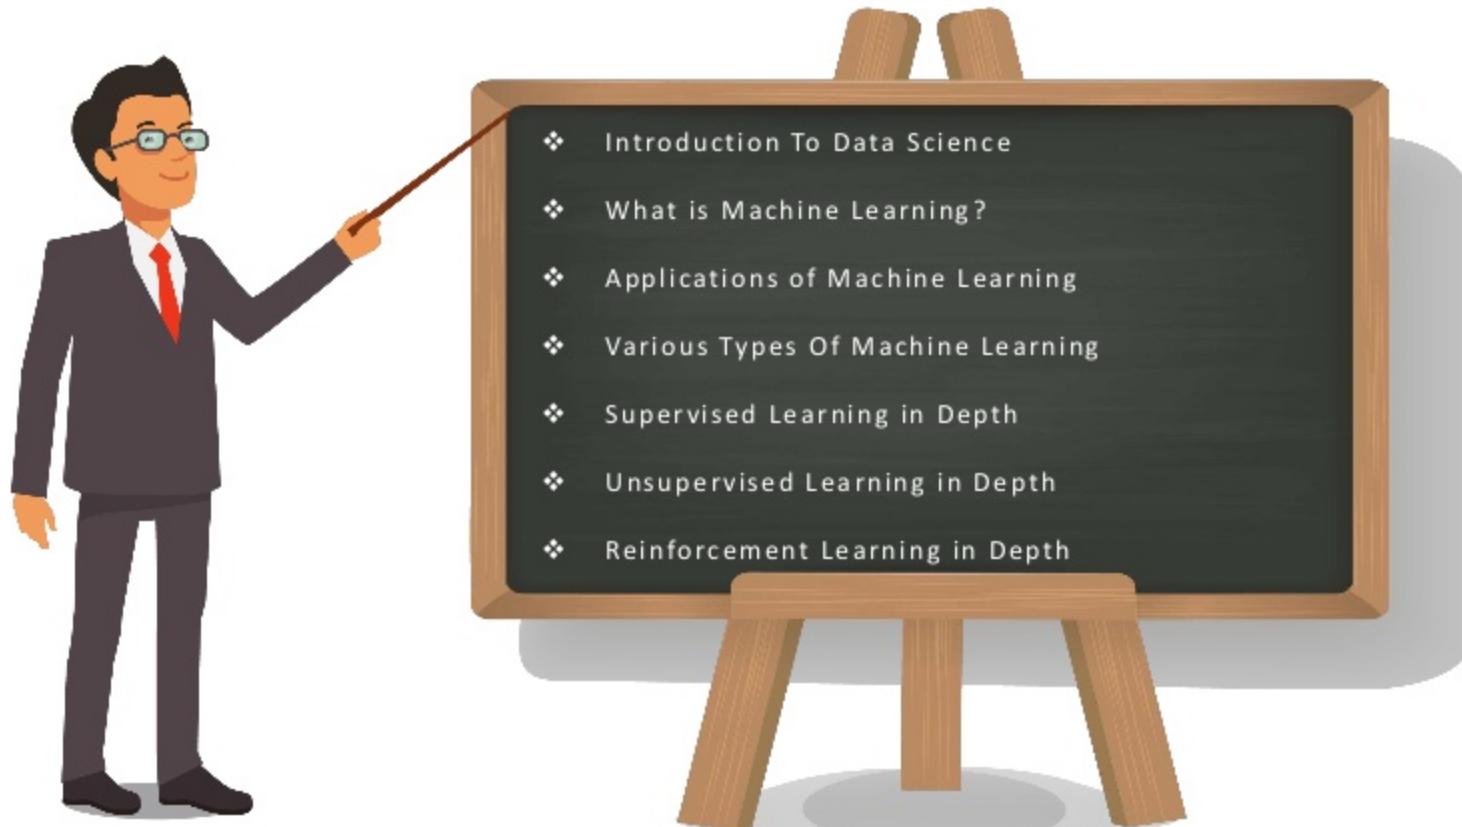


Machine Learning Full Course



Agenda

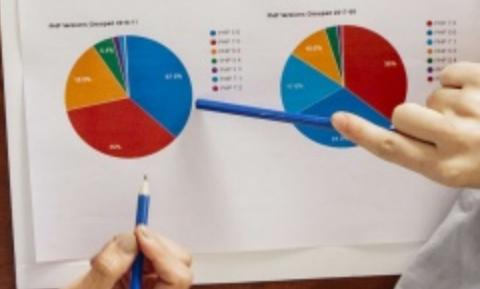


- ❖ Introduction To Data Science
- ❖ What is Machine Learning?
- ❖ Applications of Machine Learning
- ❖ Various Types Of Machine Learning
- ❖ Supervised Learning in Depth
- ❖ Unsupervised Learning in Depth
- ❖ Reinforcement Learning in Depth

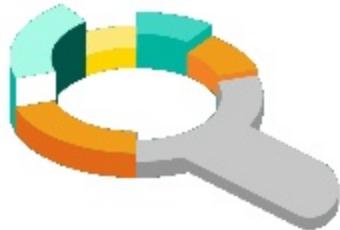
```
524184 86850  
5547321 87302  
55651 8408  
4884 887387  
587378  
54058 7308573850  
54105407 2402  
540540 088573  
57887 12098  
0405607 735856  
6020 54012  
576607 43876385  
409217 873876  
6780808 520873  
4086080 76873  
86860 87408  
87302 5703  
5408 587387  
587387 78387  
387378 38758  
54058 763  
7308573850 088573  
2402 8738  
088573 7387  
12098 38735  
735856 402736  
54012 45387  
43876385 6830540  
54308 873876  
76540  
520873 573  
76873 240857  
87408 6405  
5703 8763  
78387 0840  
38768 8578  
763 408576  
088573 840587  
8738 73830  
7387 487873  
38735 87538  
402736 763  
45387 8763  
6830540 783873  
54308 87387  
76540 673054  
573 02105766
```

```
brushTip.x-event.x  
brushTip.x-event.y-10  
brushTip.y-event.y-10  
pickColor(event.x, event.y)  
--print("moved off rect")  
brush.x:-200  
brushTip.x:W-200  
brush.y-event.y-10  
brushTip.y-event.y-10  
pickColor(W-200, event.y)  
end  
elseif event.phase=="cancelled" then  
    if dir == "left" then  
        inv = true  
        changeMouth()  
    elseif dir == "right" then  
        inv = false  
        changeMouth()  
    end  
    if self.selected == true then  
        inv = true  
        changeMouth()  
    elseif dir == "left" then  
        inv = true  
        changeMouth()  
    elseif dir == "right" then  
        inv = false  
        changeMouth()  
    end  
end
```

Which Provisioner do you use the most ? (691 responses)



What is Data Science?



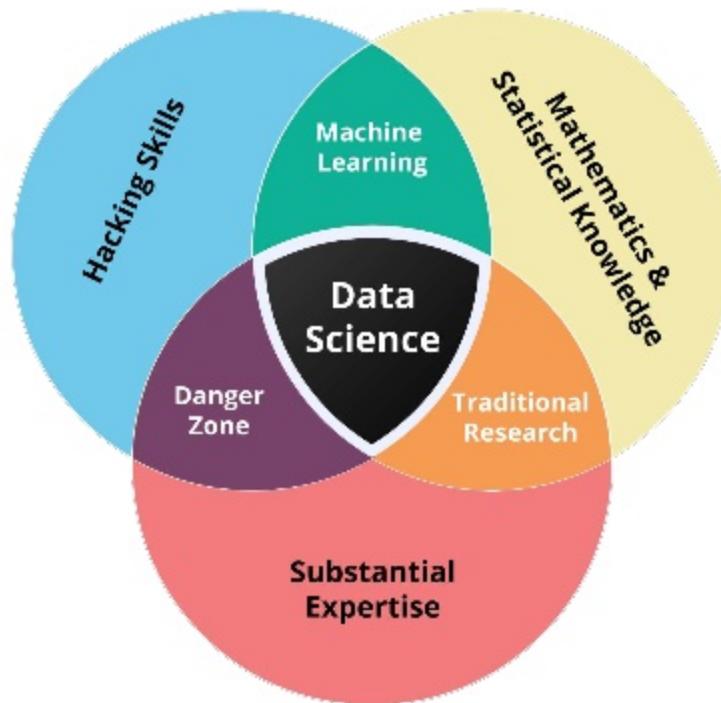
Everyday
 2.5×10^{18} Bytes



The United States alone faces a shortage of **140,000 to 190,000** people with **analytical expertise** and **1.5 million managers** and analysts with the skills to understand and make decisions based on the analysis of **big data**

What is Data Science?

Data science, also known as data-driven science, is an interdisciplinary field about scientific methods, processes, and systems to extract knowledge or insights from data in various forms, either structured or unstructured.



Data Science Peripherals



Statistics



Data Science Peripherals



Statistics



Prog Languages



Data Science Peripherals



Statistics



Prog Languages



Software



Data Science Peripherals



Statistics



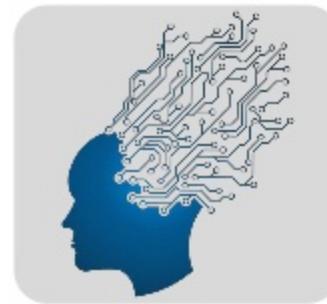
Prog Languages



Software



Machine Learning



Data Science Peripherals



Statistics



Prog Languages



Software



Machine Learning



Big Data



What is Machine Learning?

Machine Learning is a class of algorithms which is data-driven, i.e. unlike "normal" algorithms it is the data that "tells" what the "good answer" is



Getting computers to program themselves and also teaching them to make decisions using data
“Where writing software is the bottleneck, let the data do the work instead.”

Features of Machine Learning



01

It uses the data to *detect patterns* in a dataset and *adjust program actions accordingly*

It focuses on the development of computer programs that can teach themselves to grow and change when exposed to new data

02



03

It enables computers to *find hidden insights using iterative algorithms without being explicitly programmed*

Machine learning is a *method of data analysis that automates analytical model building*

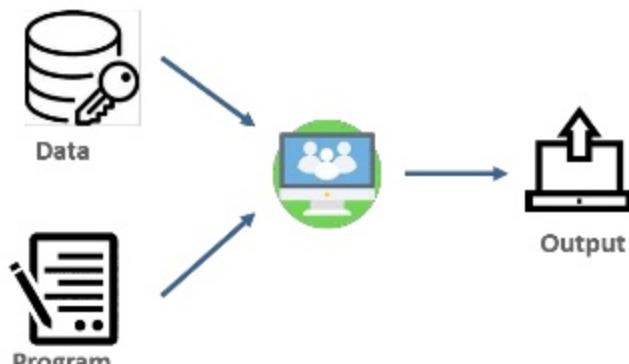
04



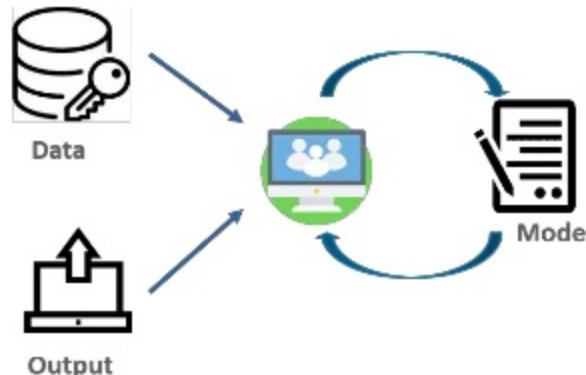
How It Works?



Traditional Programming



Machine Learning



Learn from Data

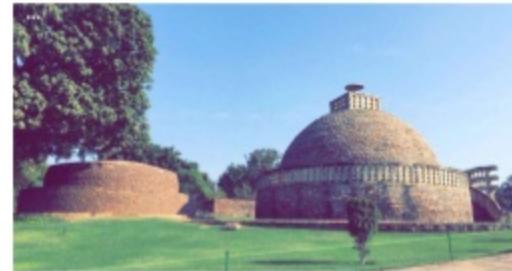
Find Hidden Insights

Train and Grow

Applications of Machine Learning

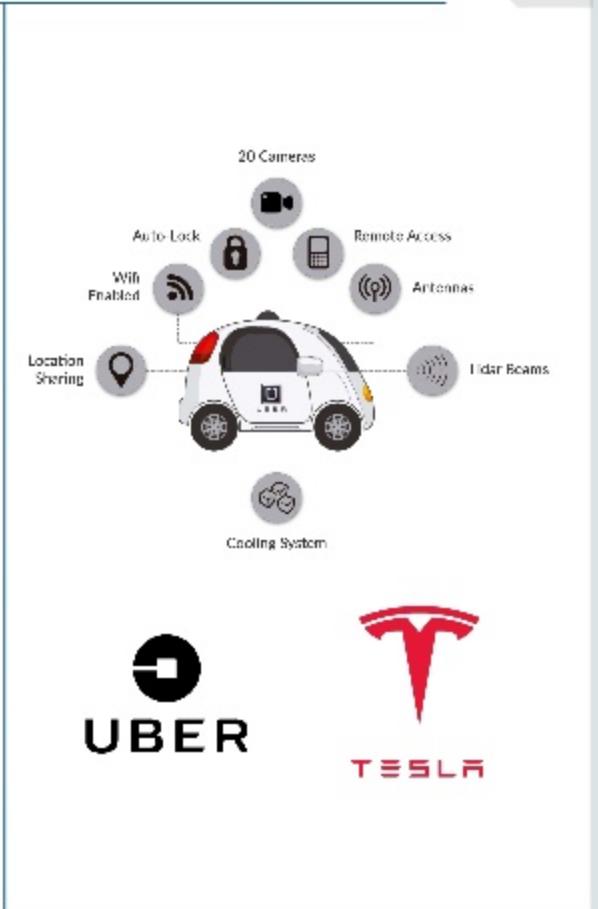
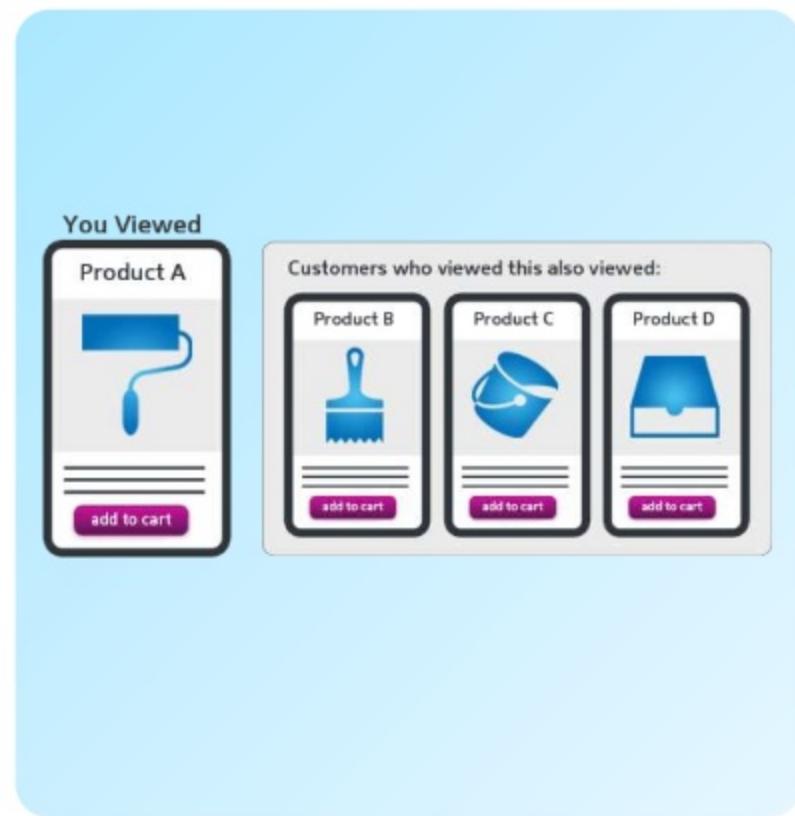


Gary Chavez added a photo you might ...
be in.
about a minute ago ·

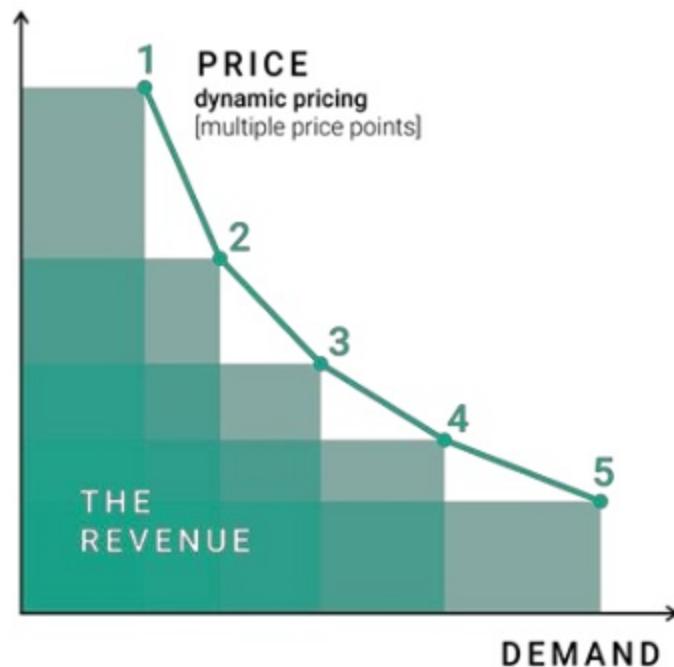


UBER

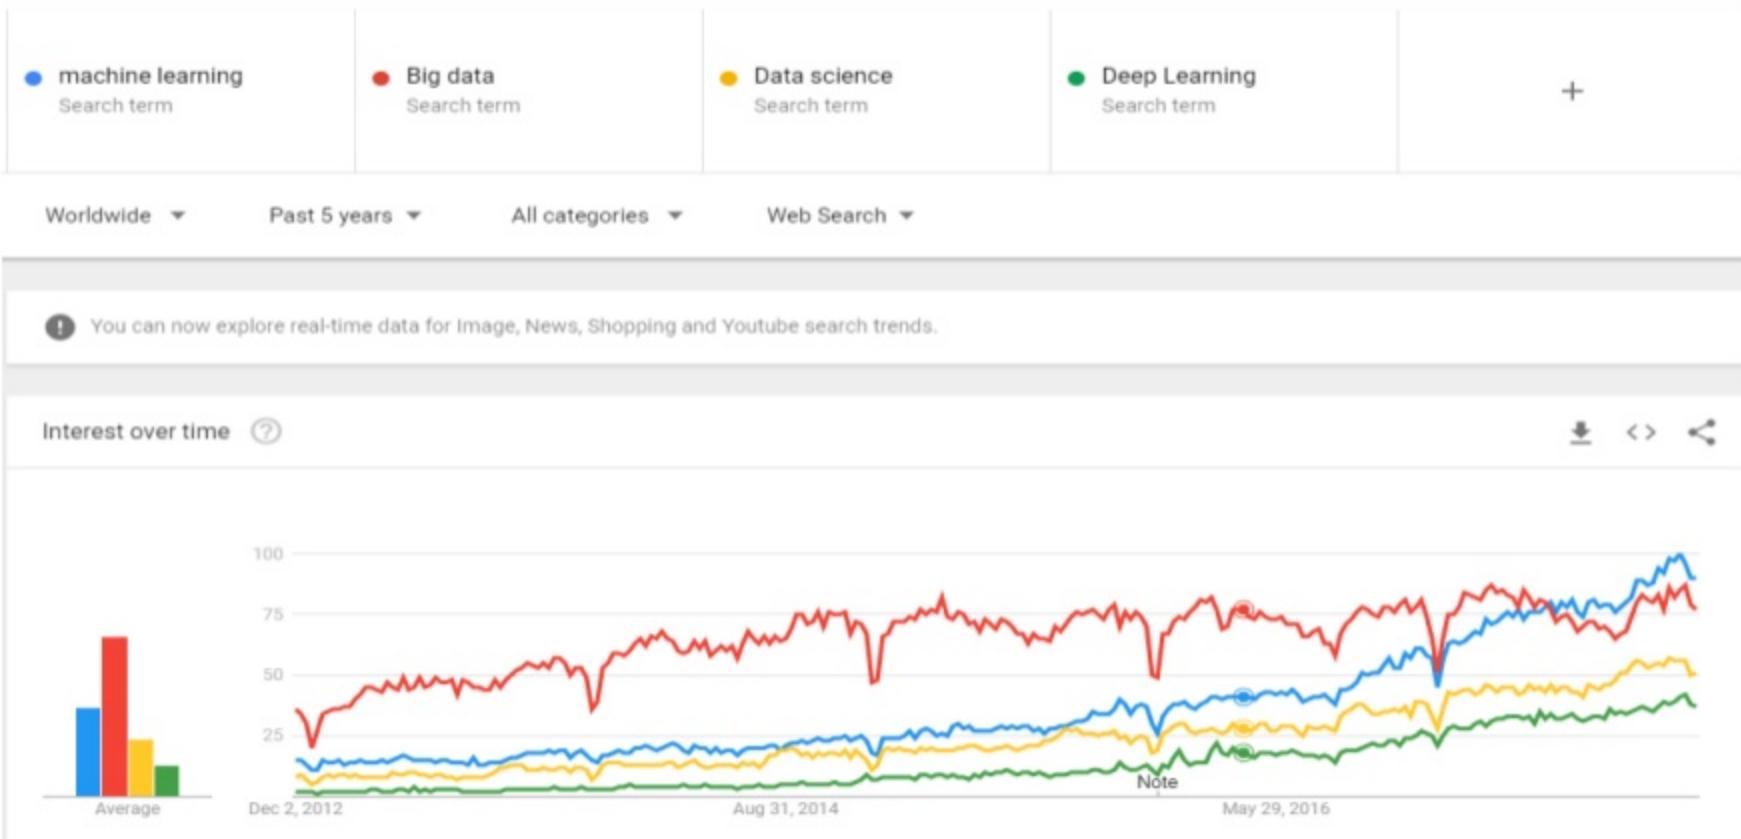
Applications of Machine Learning



Applications of Machine Learning



Market Trend: Machine Learning

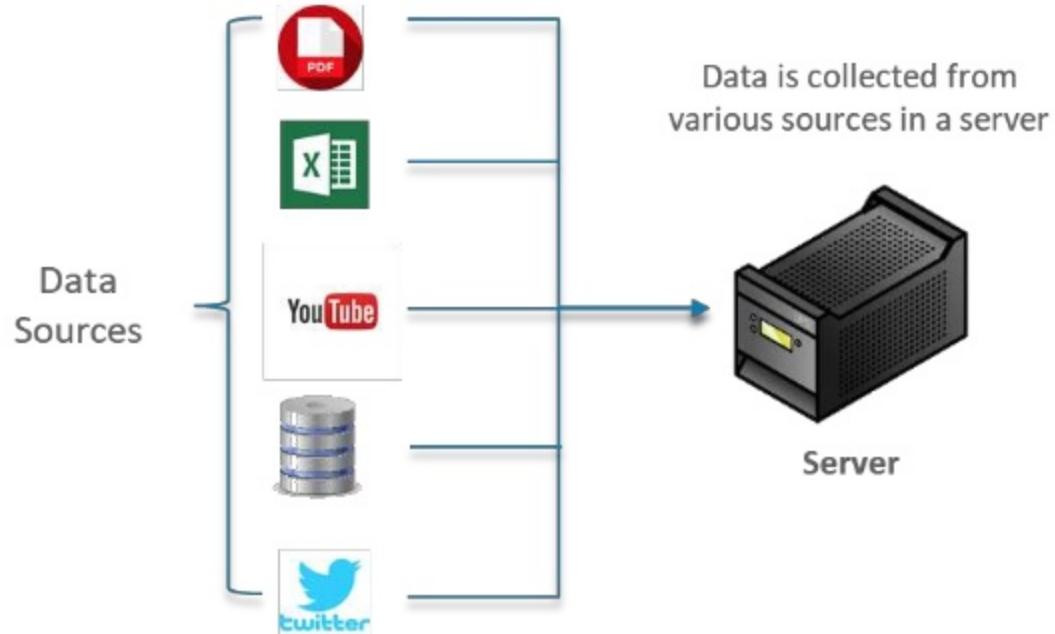


Machine Learning Life Cycle



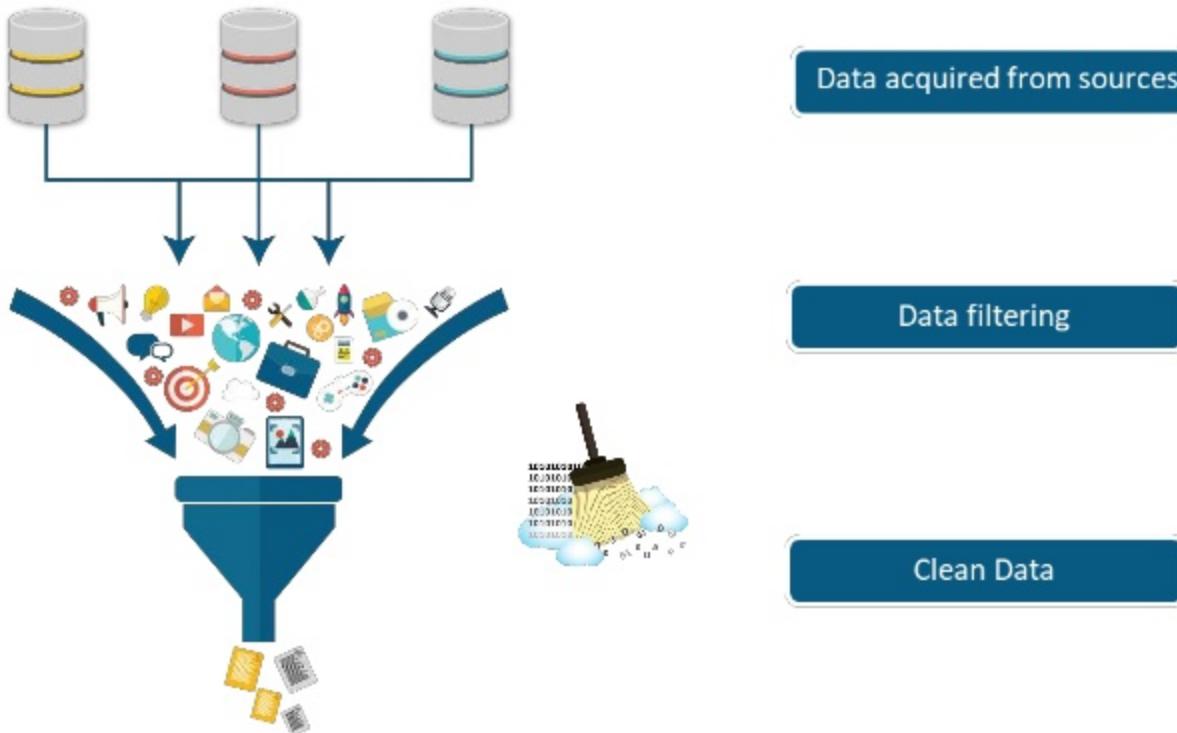
Step 1: Collecting Data

- 1
- 2
- 3
- 4
- 5
- 6



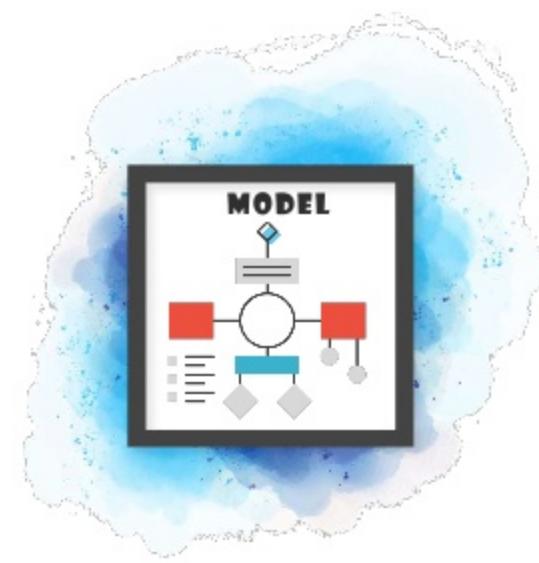
Data Wrangling

- 1
- 2
- 3
- 4
- 5
- 6



Analyse Data

- 1
- 2
- 3
- 4
- 5
- 6



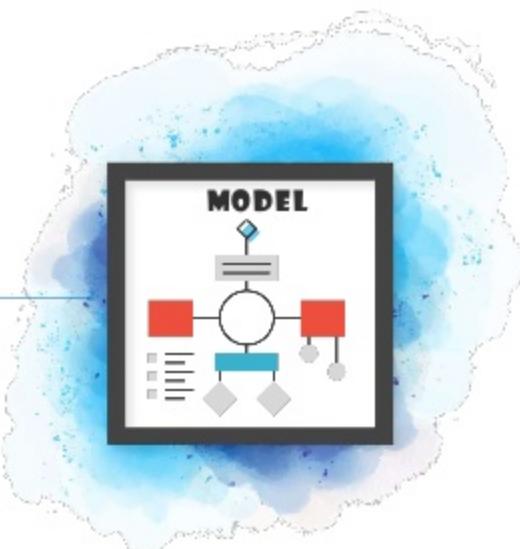
Train Algorithm

- 1
- 2
- 3
- 4
- 5
- 6

Player Name	Date birth
Manuel TORNER	1984-06-22
Oscar REYES	1986-04-13
Marc ALEXANDRE	1985-07-28
David CARRASCO	1984-05-23
Wenceslao MOLINA	1983-05-22
David GARCIA	1985-06-11
Alberto VILLALBA	1985-04-18
Adrián ALVAREZ	1984-07-10
Antonio MELLO	1985-08-09
Oscar GARCIA	1986-03-11
Miguel PÉREZ	1985-06-01
Alex BELTRAN	1986-02-02
Manel RIBERA	1985-03-01
Alejo SOTO	1985-07-20
Mercedes MUÑOZ	1984-05-19
Sergio GARCIA	1985-02-19
Jesús AMOROS	1984-05-01

Training set

Medication	Days since
Hypoglycemic	14
Antihypertensive	30
Cholesterol	30
Diabetes	14
Arthritis	14
Hypertension	14
Urinary tract infection	30
Heart disease	30
Stroke	30
Parkinson's	14
Back pain	14
Osteoporosis	30
Peptic ulcer	30
Rheumatoid arthritis	14
Hormones	14
Arthritis	14
Fever	14
Urinary tract infection	30
Glaucoma	30
Varicose veins	30
Thyroid condition	30
Aspirin	14



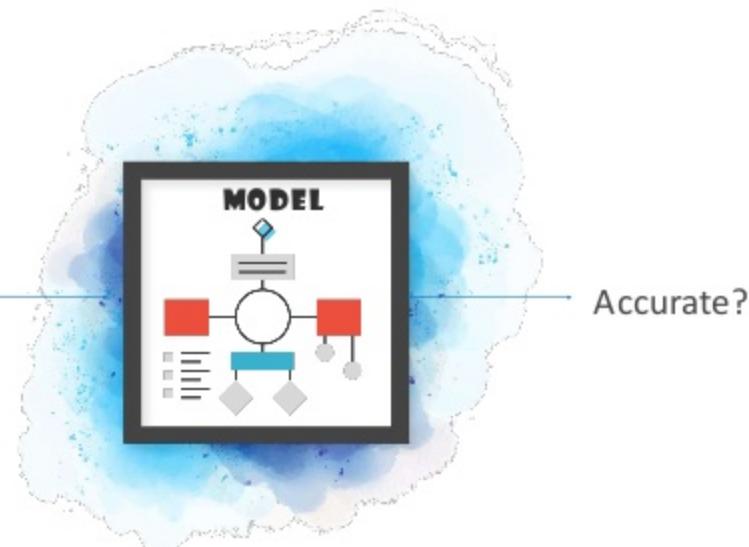
Test Algorithm

- 1
- 2
- 3
- 4
- 5
- 6

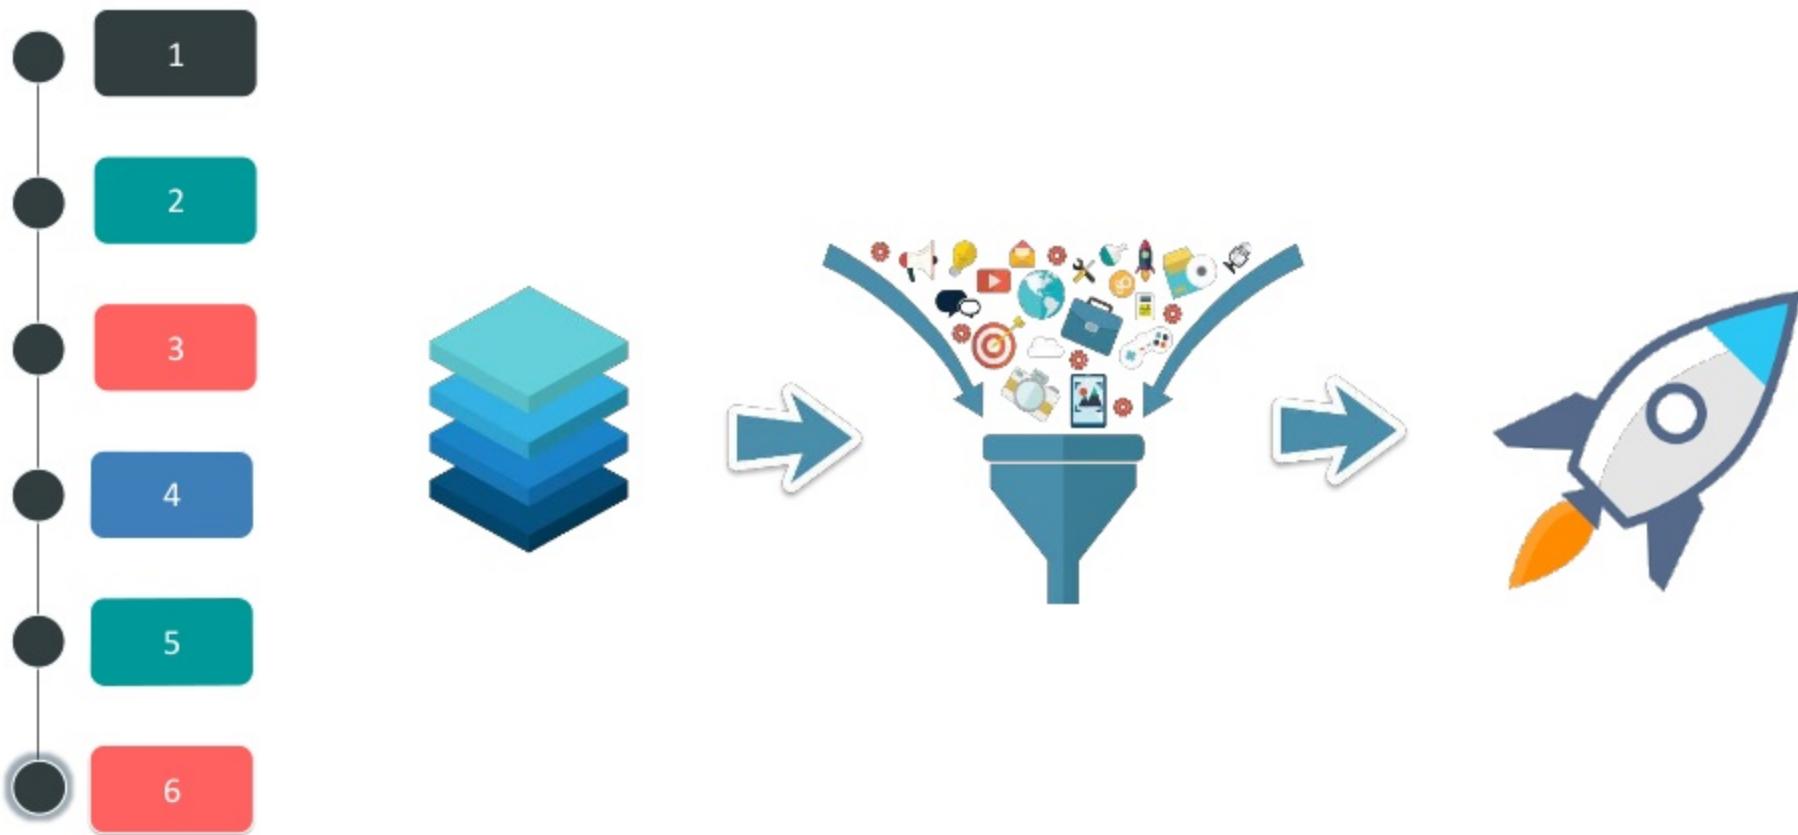
Player Name	Death Name
Adam Tappert	LICHEN Head
Oscar WEPFELD	LICHEN Head
Mario MUNIZ	LICHEN Head
Mike CARRIED	LICHEN Head
Wesley LEE	LICHEN Head
Robert GRIFF	LICHEN Head
andrea FALCON	LICHEN Head
Alvaro MELO	LICHEN Head
Domenic MELLO	LICHEN Head
Diego MELLO	LICHEN Head
Henrique MELLO	LICHEN Head
Miguel PERIN	LICHEN Head
Alex REZENDE CRISTOVES	LICHEN Head
Paulo RIBEIRO	LICHEN Head
Juliano RIBEIRO	LICHEN Head
Hernani RIBEIRO	LICHEN Head
Edson RIBEIRO	LICHEN Head
Antonio RIBEIRO	LICHEN Head
Adriano RIBEIRO	LICHEN Head

Test set

Macrophage	CD14+
Neutrophil	10
Lympohoid	10
Leukocyte	10
Monocyte	10
Neutrophil	10
Monocyte	10
Leukocyte	10
Neutrophil	10
Monocyte	10



Operation and Optimization



Important Python Libraries

Seaborn

Focused on the visual of statistical models which include heat maps and depict the overall distributions

Matplotlib

It enables you to make- Bar charts, Scatter plots, Line Charts, Histograms, Pie charts, Contour plots, Quiver plots

Scikit-Learn

Simple and efficient for data mining and data analysis, Built on NumPy and matplotlib, Open source

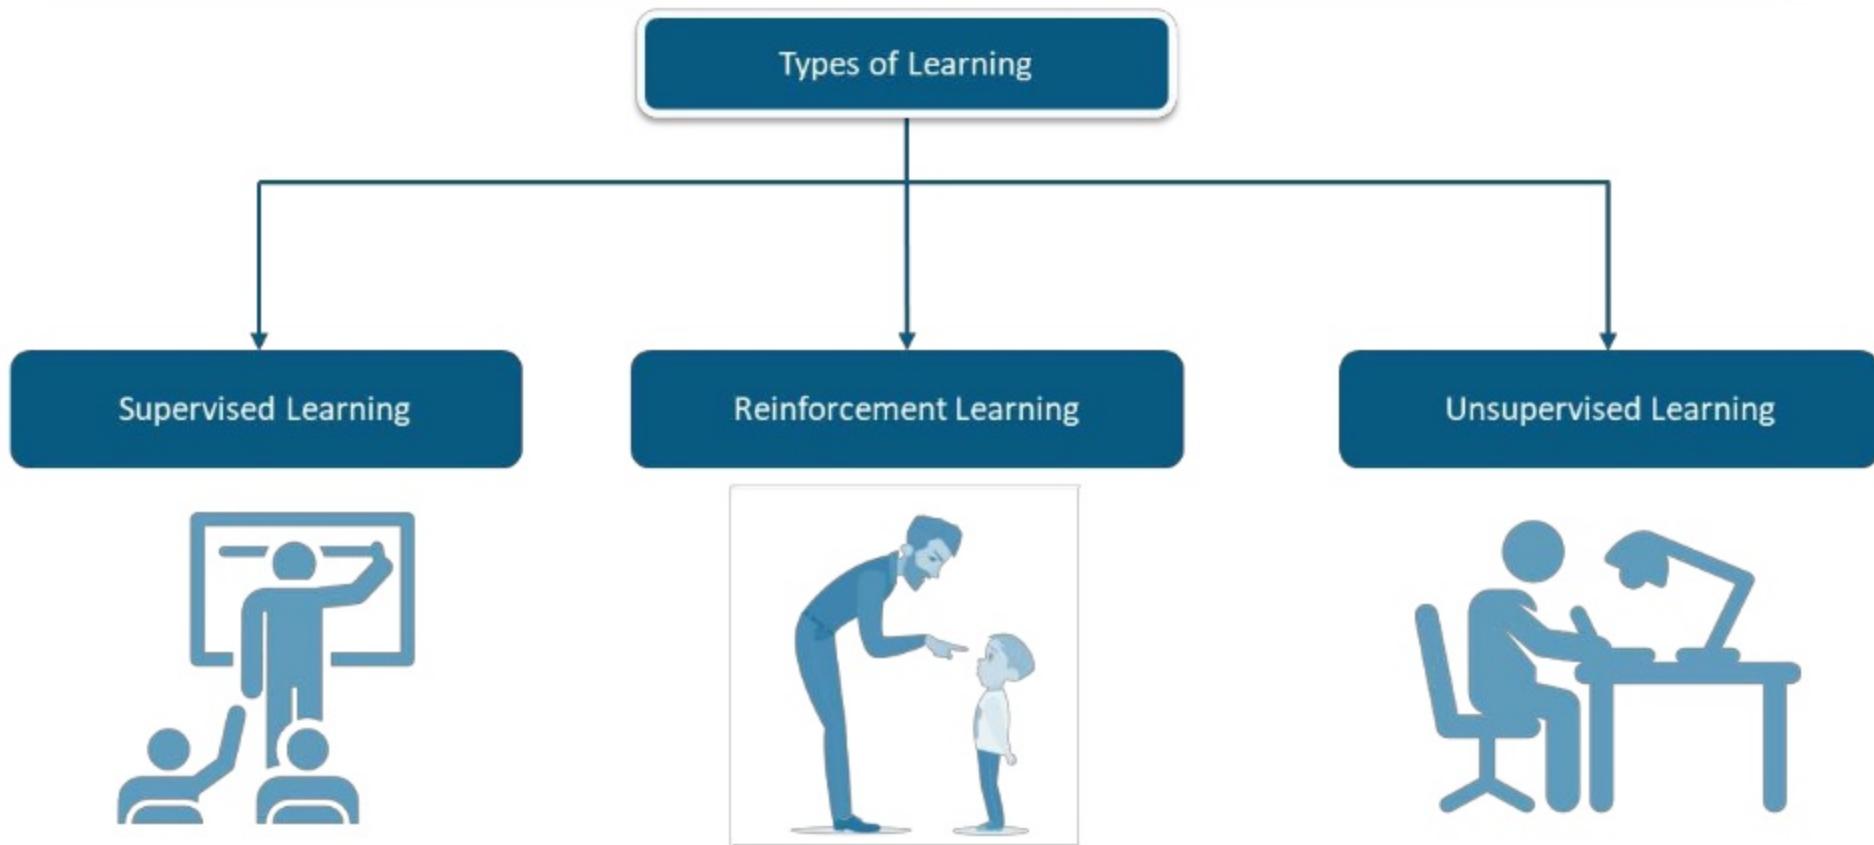
Pandas

Perfect tool for data wrangling, designed for quick and easy data manipulation, aggregation, and visualization

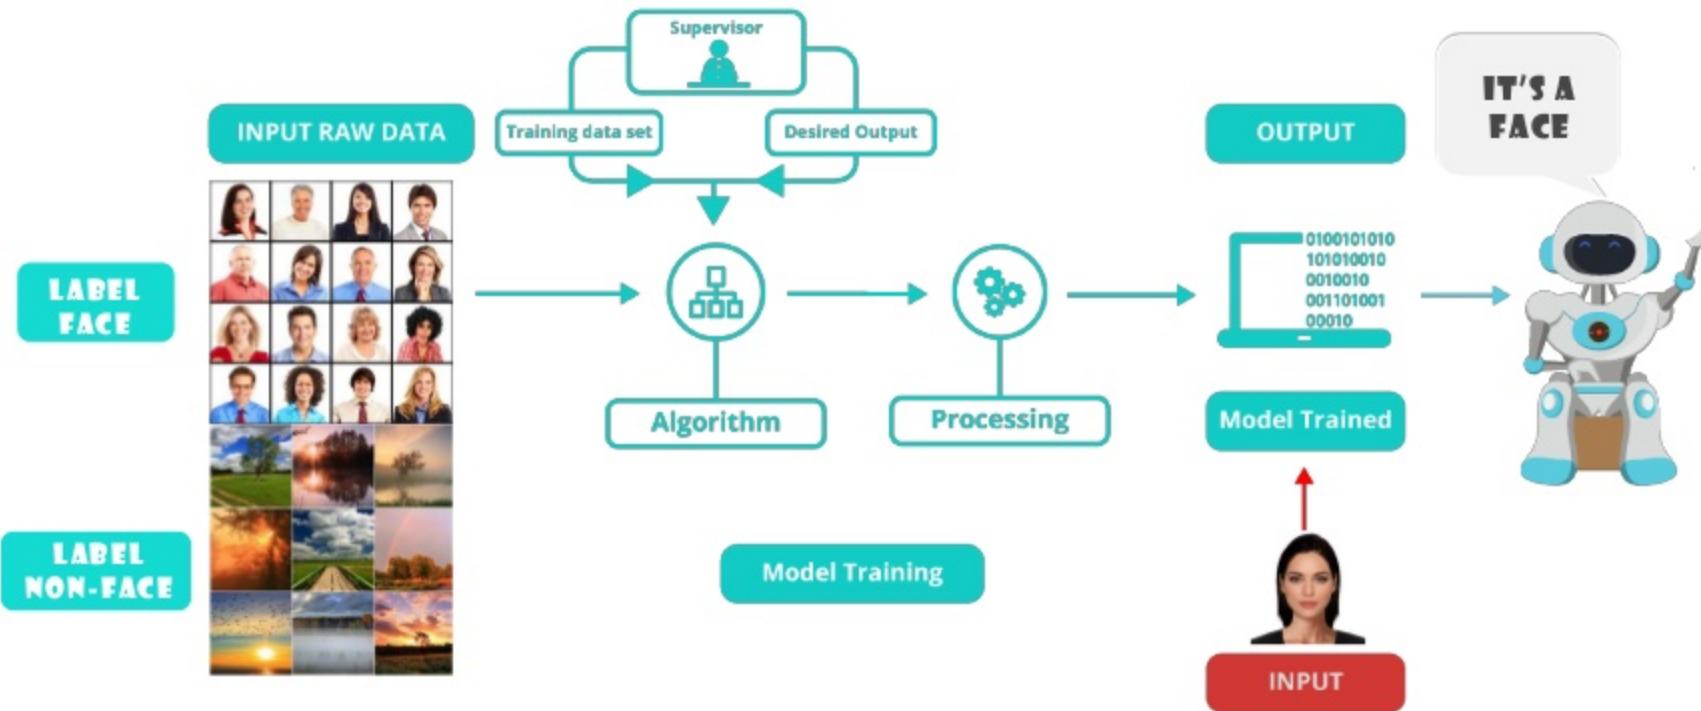
Numpy

Stands for Numerical Python, provides an abundance of useful features for operations on n-arrays and matrices in Python

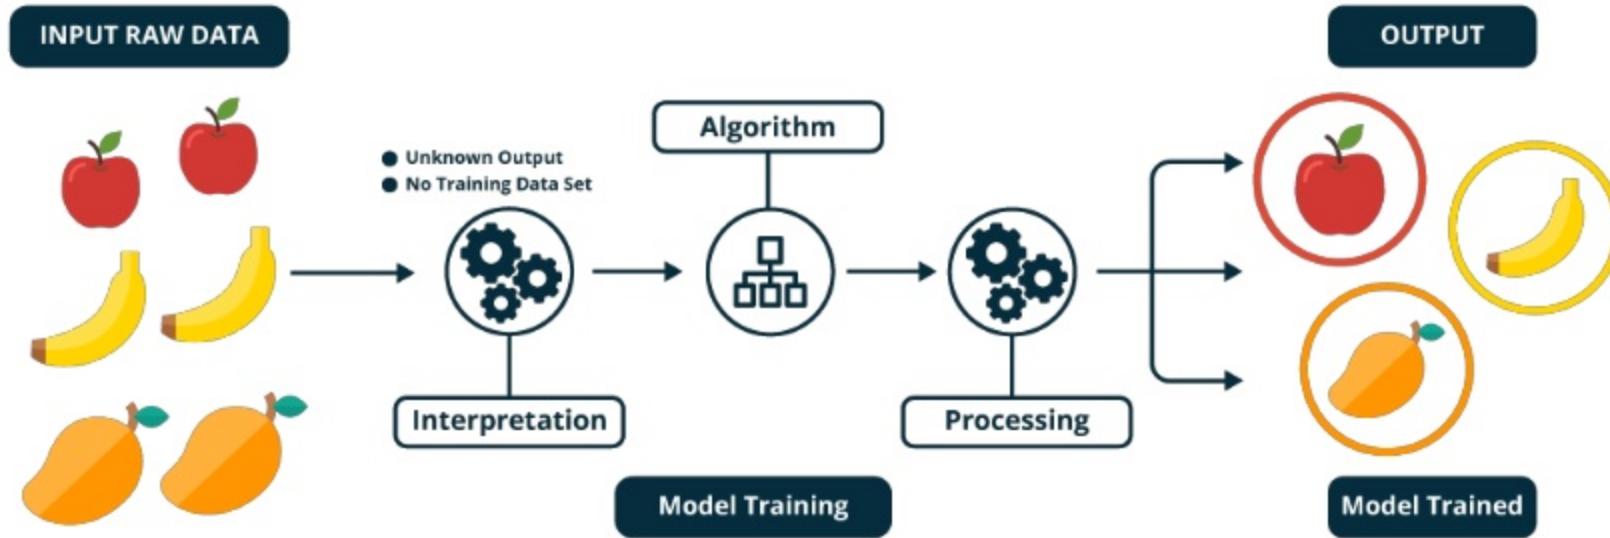
Types of Machine Learning



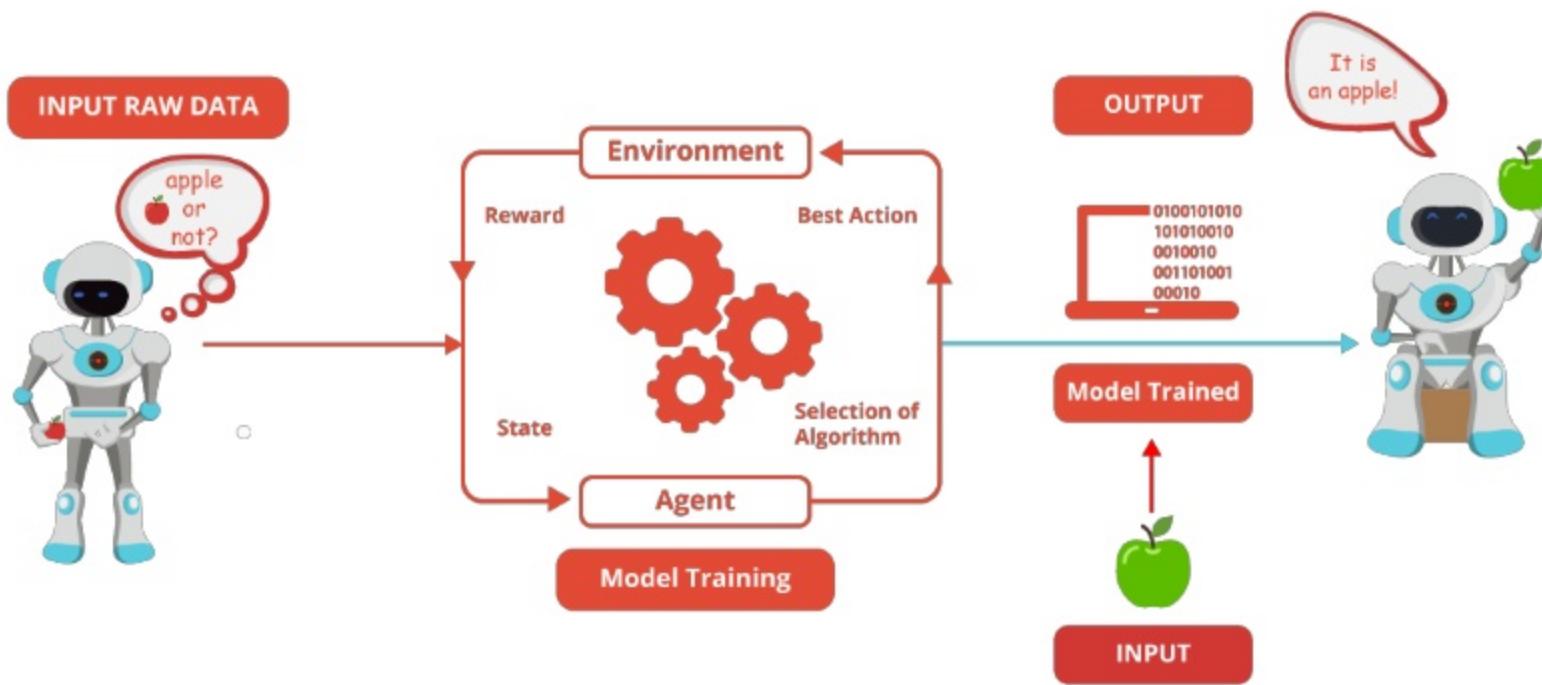
Supervised Learning



Unsupervised Learning

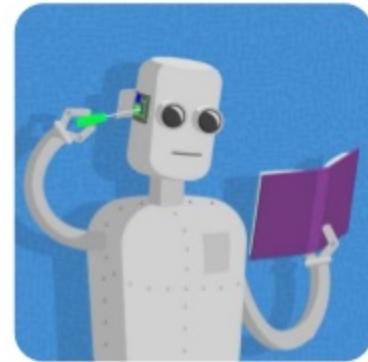


Reinforcement Learning



Supervised Learning

Supervised learning is where you have input variables (x) and an output variable (Y) and you use an algorithm to learn the mapping function from the input to the output



It is called Supervised Learning because the process of an algorithm learning from the training dataset can be thought as a teacher supervising the learning process

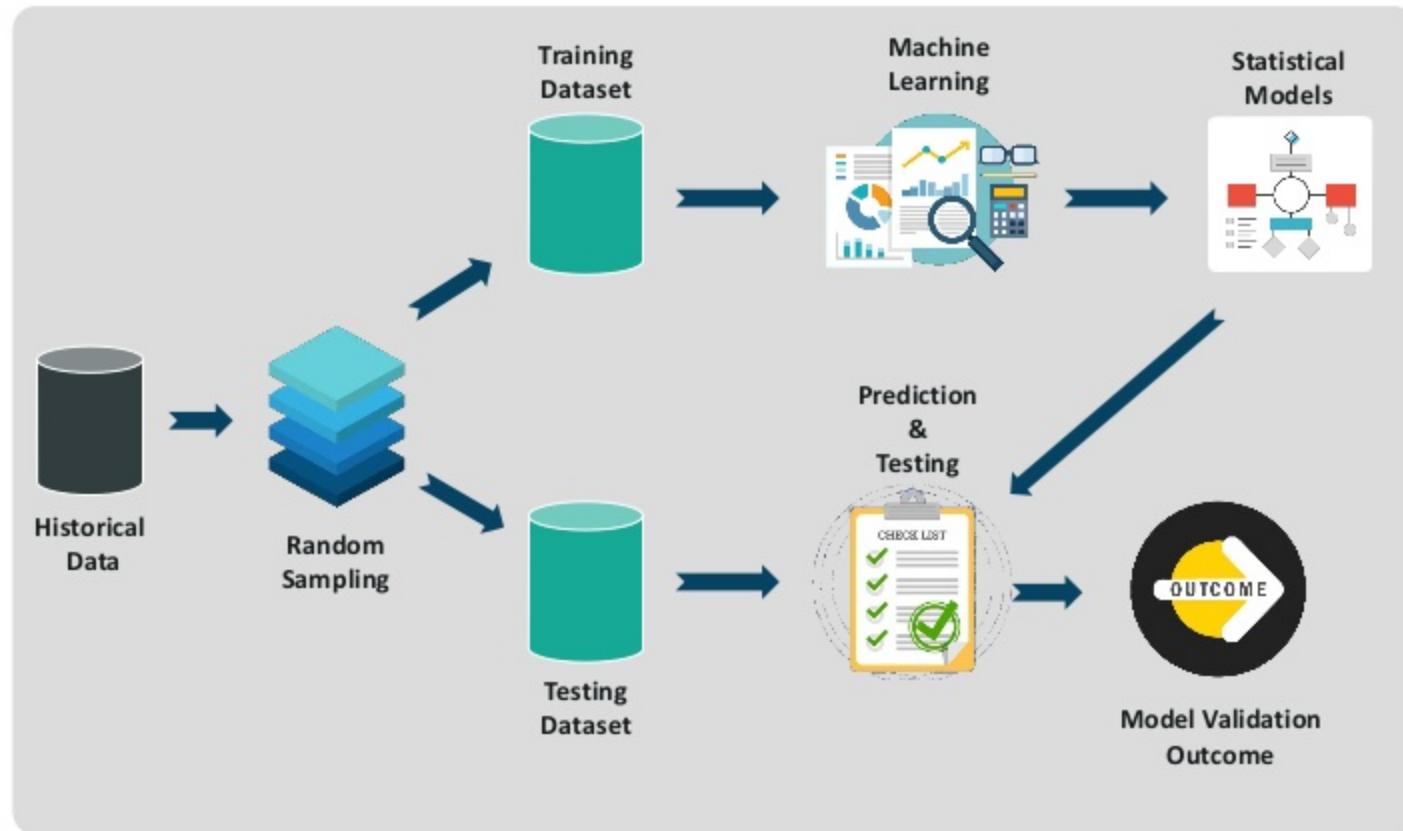
Supervised Learning



Training and Testing



Prediction



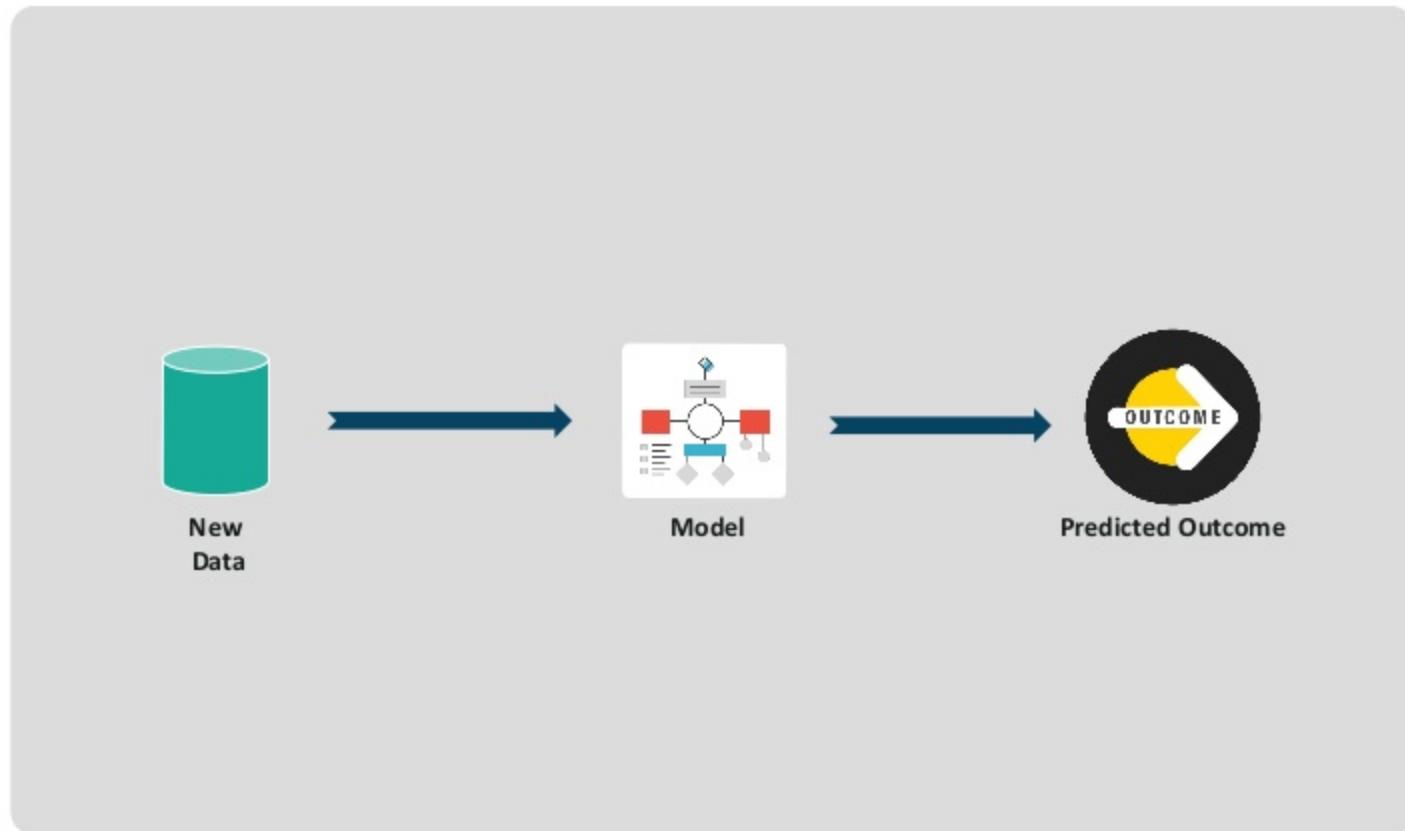
Supervised Learning



Training and Testing



Prediction



Supervised Learning Algorithms



Linear Regression



Logistic Regression



Decision Tree



Random Forest

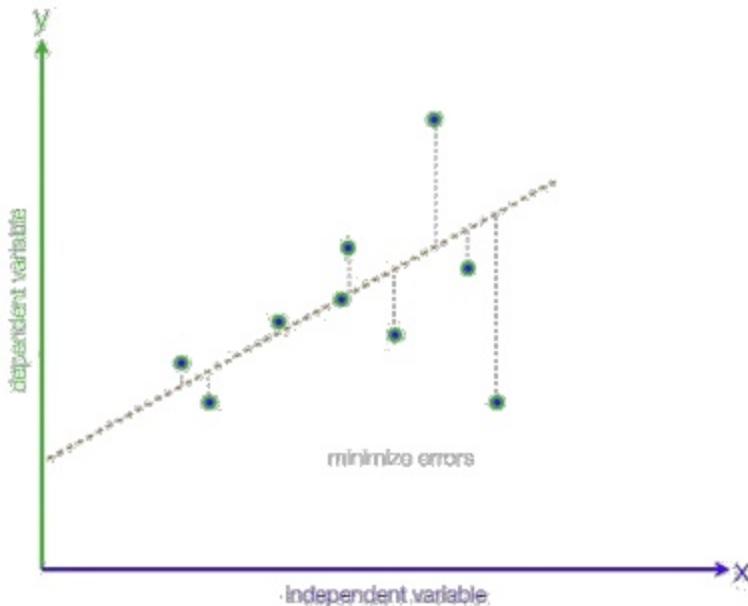


Naïve Bayes Classifier

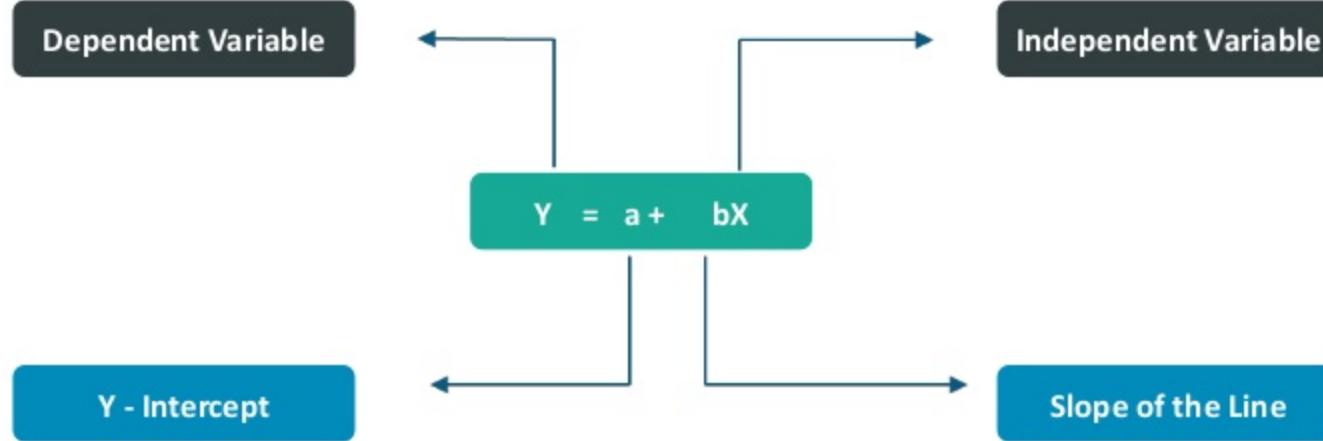
Linear Regression

Linear Regression Analysis is a powerful technique used for predicting the unknown value of a variable (**Dependent Variable**) from the known value of another variables (**Independent Variable**)

- A **Dependent Variable(DV)** is the variable to be predicted or explained in a regression model
- An **Independent Variable(IDV)** is the variable related to the dependent variable in a regression equation

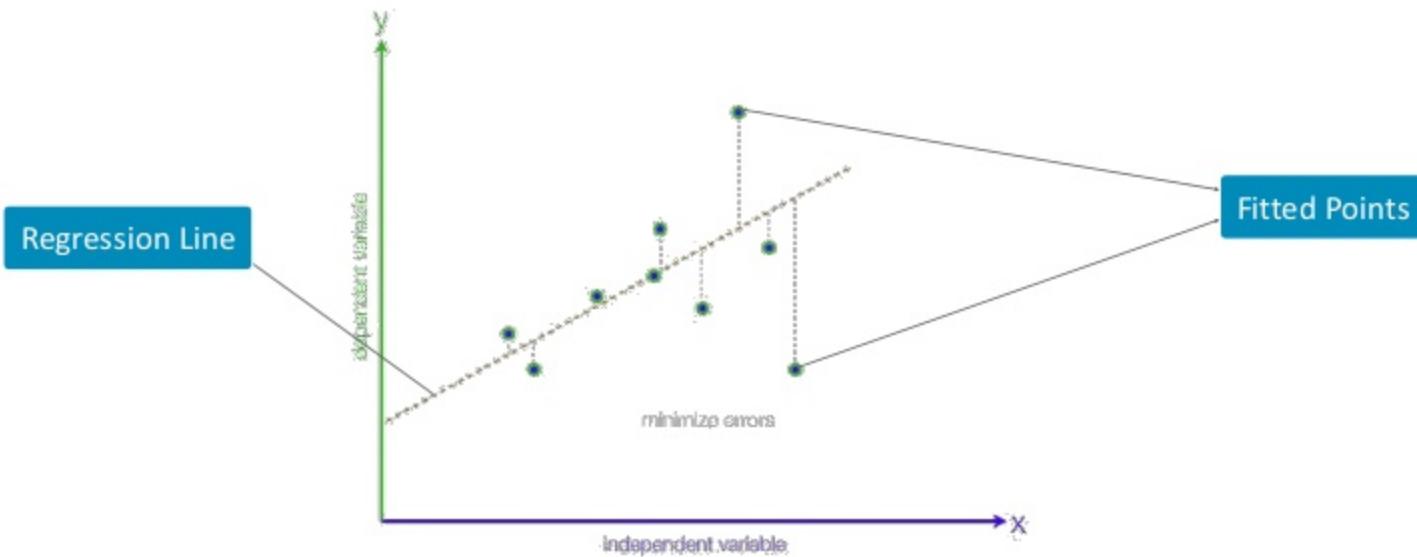


Simple Linear Regression



Regression Line

Linear Regression Analysis is a powerful technique used for predicting the unknown value of a variable (**Dependent Variable**) from The regression line is simply a single line that best fits the data
(In terms of having the smallest overall distance from the line to the points)



Demo

Real Estate Company Use Case



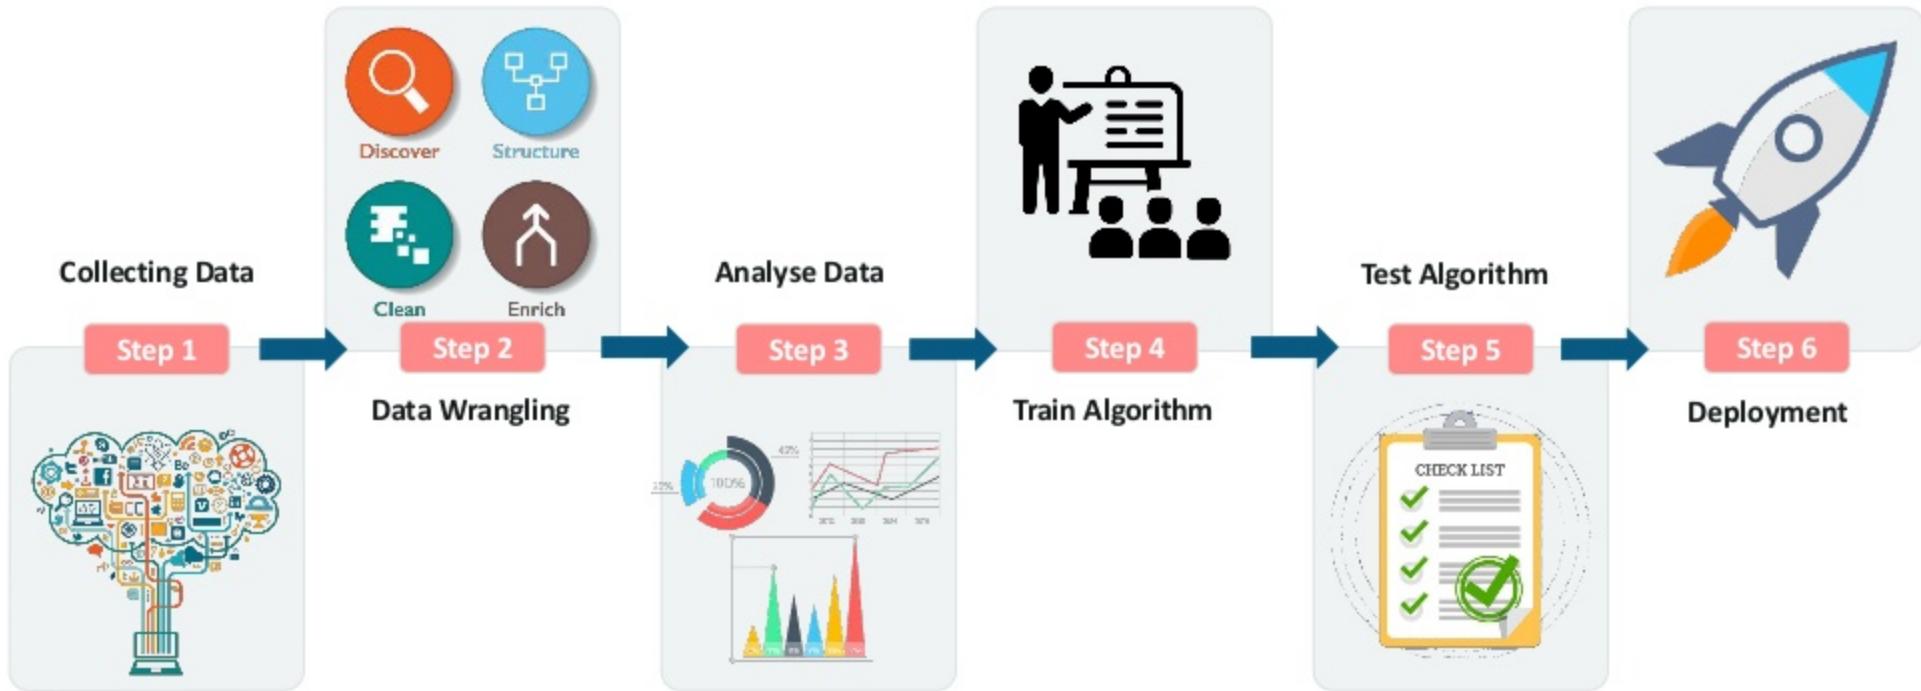
Hi I am John, I need some baseline for pricing my Villas and Independent Houses



Dataset Description

Column	Description
CRIM	per capita crime rate by town
ZN	proportion of residential land zoned for lots over 25,000 sq.ft.
INDUS	proportion of non-retail business acres per town.
CHAS	Charles River dummy variable (1 if tract bounds river; 0 otherwise)
NOX	nitric oxides concentration (parts per 10 million)
RM	average number of rooms per dwelling
AGE	proportion of owner-occupied units built prior to 1940
DIS	weighted distances to five Boston employment centres
RAD	index of accessibility to radial highways
TAX	full-value property-tax rate per \$10,000
PTRATIO	pupil-teacher ratio by town
B	$1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town
LSTAT	% lower status of the population
MEDV	Median value of owner-occupied homes in \$1000's

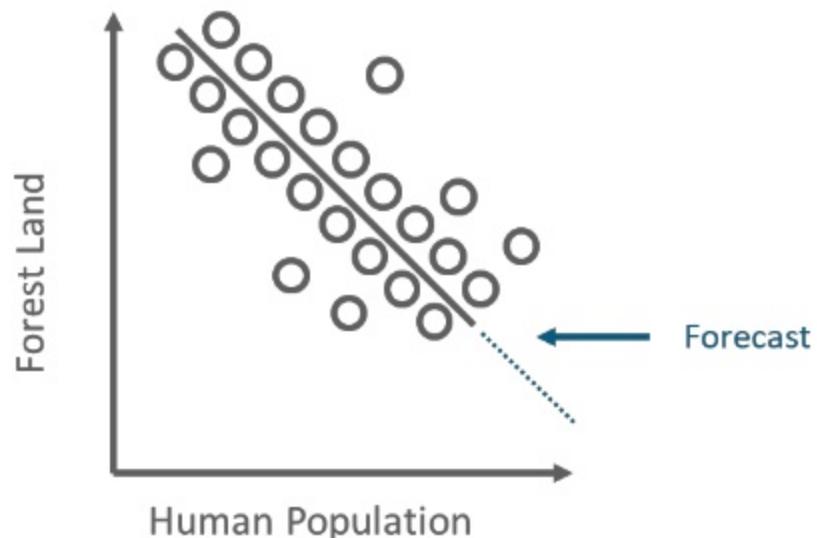
Steps



Model Fitting

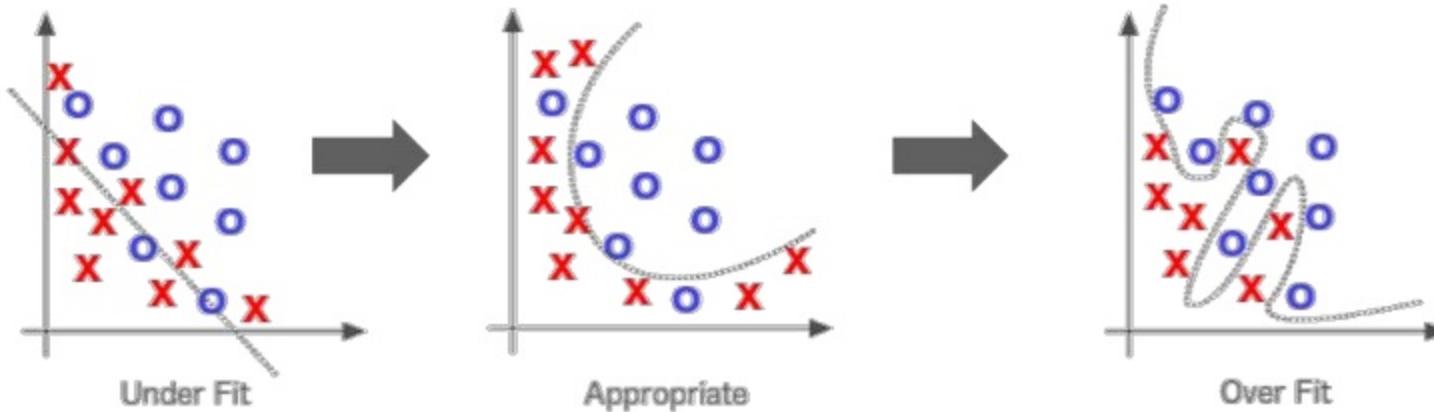
Fitting a model means that you're making your algorithm learn the relationship between predictors and outcome so that you can predict the future values of the outcome .

So the best fitted model has a specific set of parameters which best defines the problem at hand



Types of Fitting

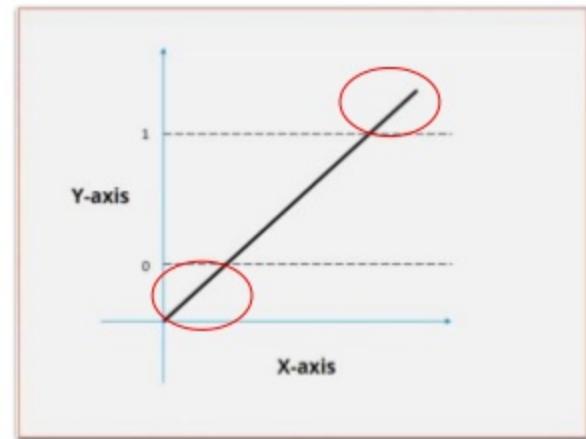
Machine Learning algorithms first attempt to solve the problem of under-fitting; that is, of taking a line that does not approximate the data well, and making it to approximate the data better.



Need For Logistic Regression



Here, the best fit line in linear regression is going below 0 and above 1



WHO WILL WIN ?

What is Logistic Regression?

Logistic Regression is a statistical method for analysing a dataset in which there are one or more independent variables that determine an outcome.
Outcome is a binary class type.

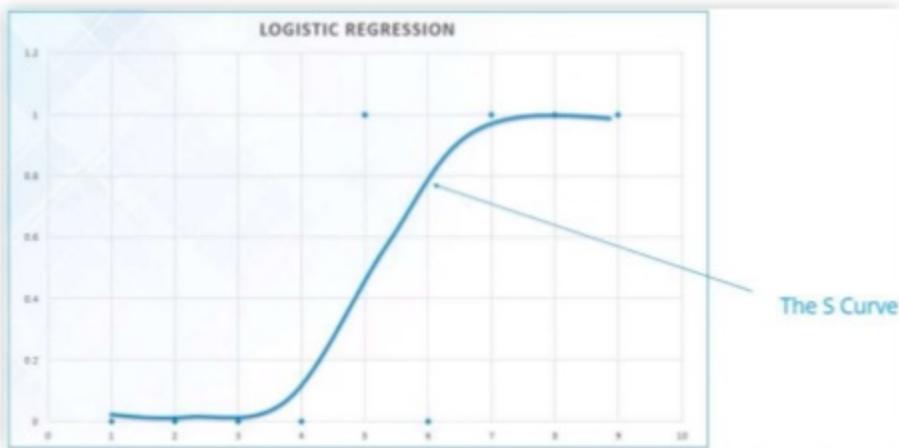


The outcome(result)
will be binary(0/1)

- 0- If malignant
- 1- If benign

What is Logistic Regression?

The Logistic Regression Curve is called as “**Sigmoid Curve**”, also known as **S-Curve**

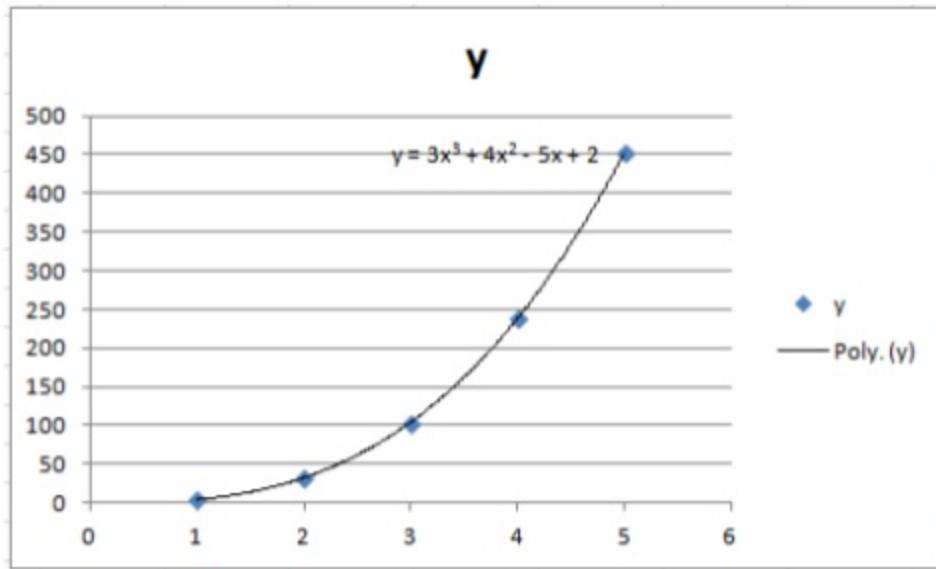


Based on the threshold value set, we decide the output from the function

What is Polynomial Regression?

When we have non linear data, which can't be predicted with a linear model. We switch to Polynomial Regression.

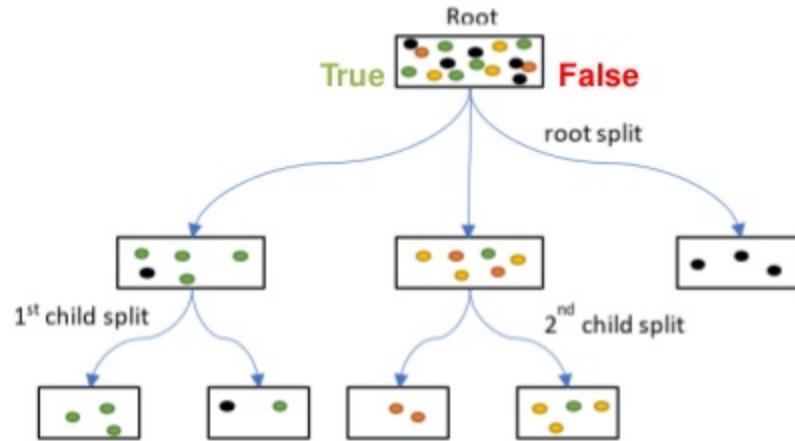
Such a scenario is shown in the below graph



What is a Decision Tree?

A decision tree is a tree-like structure in which internal node represents test on an attribute

- Each branch represents outcome of test and each leaf node represents class label (decision taken after computing all attributes)
- A path from root to leaf represents classification rules.



Building a Decision Tree



Day	Outlook	Humidity	Wind
D1	Sunny	High	Weak
D2	Sunny	High	Strong
D8	Sunny	High	Weak
D9	Sunny	Normal	Weak
D11	Sunny	Normal	Strong

2 Yes / 3 No

Split further

Day	Outlook	Humidity	Wind
D3	Overcast	High	Weak
D7	Overcast	Normal	Strong
D12	Overcast	High	Strong
D13	Overcast	Normal	Weak

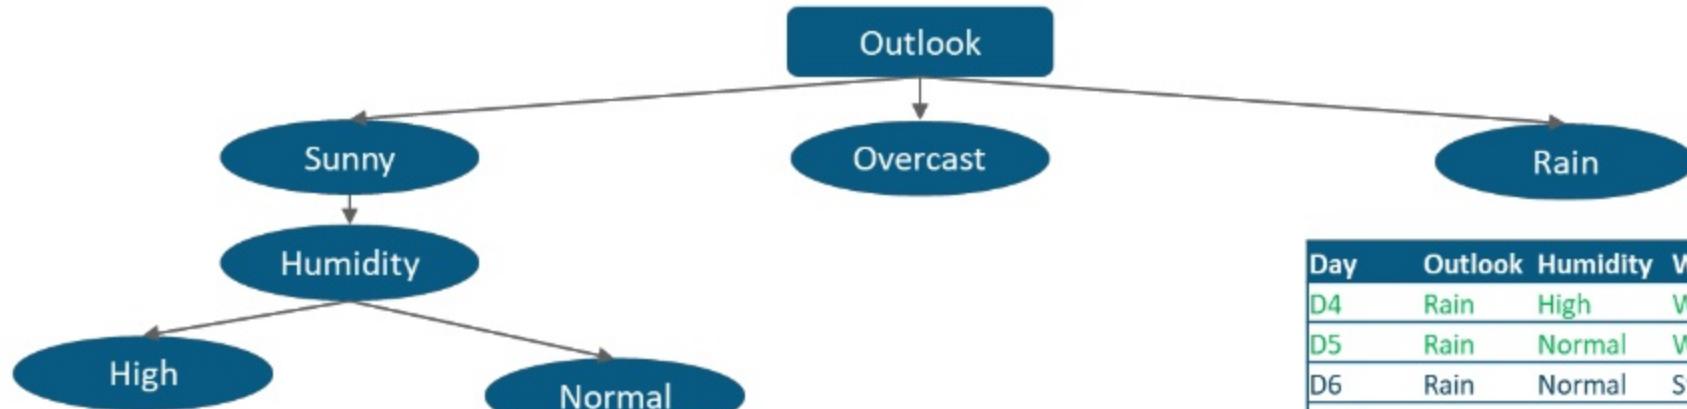
Pure subset
Yes(will play)

Day	Outlook	Humidity	Wind
D4	Rain	High	Weak
D5	Rain	Normal	Weak
D6	Rain	Normal	Strong
D10	Rain	Normal	Weak
D14	Rain	High	Strong

3 Yes / 2 No

Split further

Building a Decision Tree



Day	Humidity	Wind
D1	High	Weak
D2	High	Strong
D8	High	Weak

Pure subset

NO(will not play)

Day	Humidity	Wind
D9	Normal	Weak
D11	Normal	Strong

Pure subset

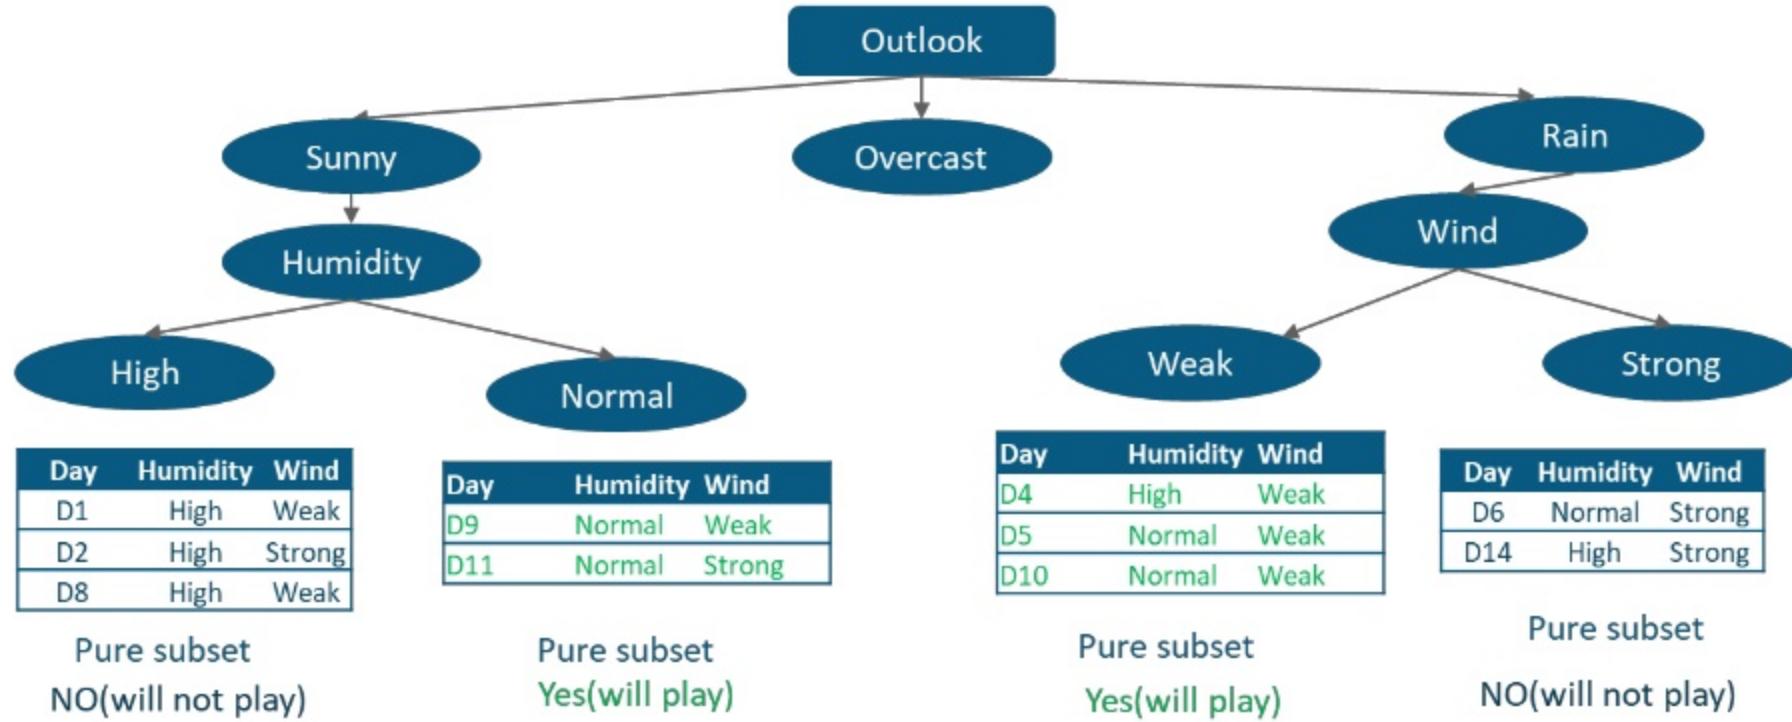
Yes(will play)

Day	Outlook	Humidity	Wind
D4	Rain	High	Weak
D5	Rain	Normal	Weak
D6	Rain	Normal	Strong
D10	Rain	Normal	Weak
D14	Rain	High	Strong

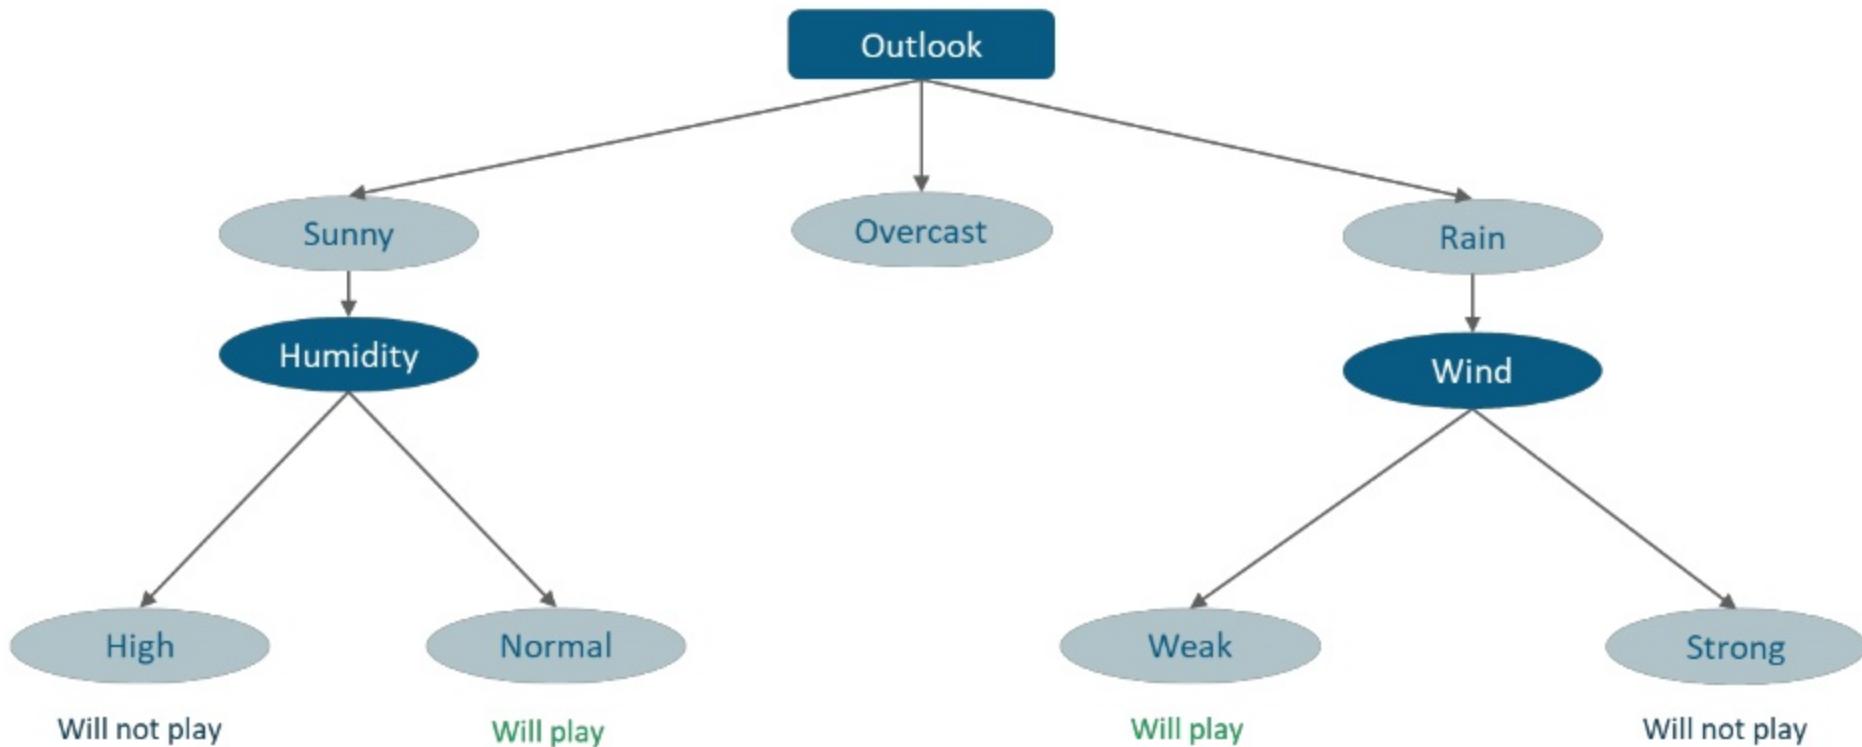
3 Yes / 2 No

Split further

Building a Decision Tree



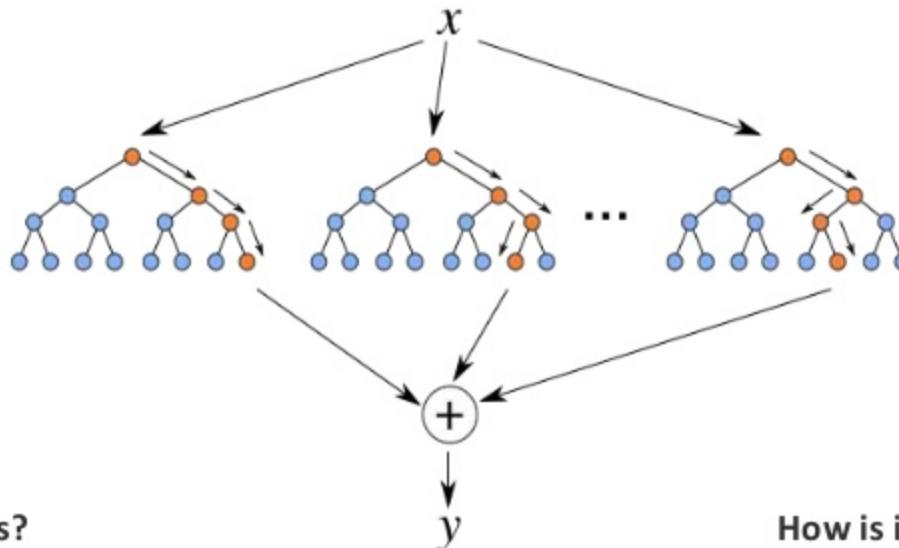
Building a Decision Tree



Demo

What is Random Forest?

Random Forest is an ensemble classifier made using many Decision tree models



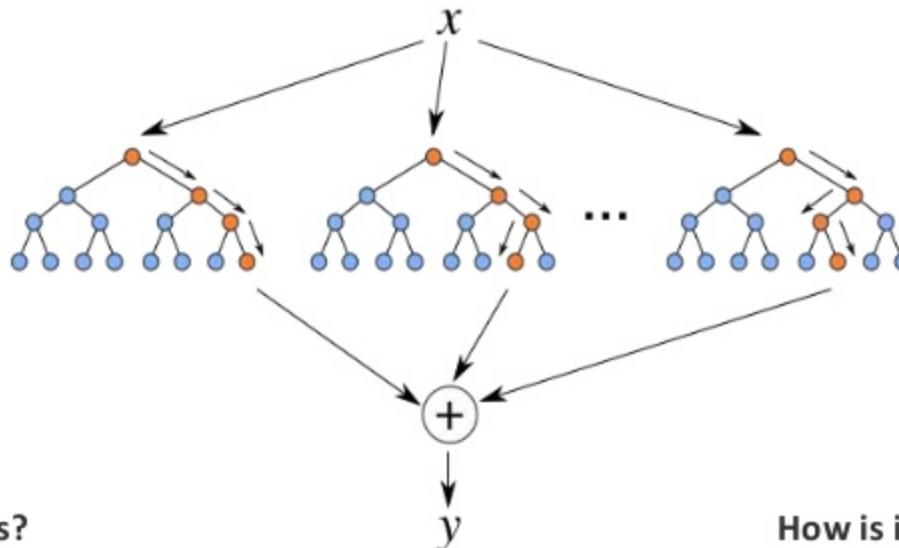
What are Ensemble models?

How is it better from Decision Trees ?

Demo

What is Naïve Bayes?

Random Forest is an ensemble classifier made using many Decision tree models



What are Ensemble models?

How is it better from Random Forest?

What is Naïve Bayes ?



Naïve Bayes is a simple but surprisingly powerful algorithm for predictive modeling.



Classification Technique

Bayes' Theorem

Given a hypothesis H and evidence E , Bayes' theorem states that the relationship between the probability of the hypothesis before getting the evidence $P(H)$ and the probability of the hypothesis after getting the evidence $P(H|E)$ is

$$P(H|E) = \frac{P(E|H).P(H)}{P(E)}$$



Bayes' Theorem Example



Bayes' Theorem Example



$$P(\text{King}) = 4/52 = 1/13$$

Bayes' Theorem Example



$$P(\text{King}) = 4/52 = 1/13$$

$$P(\text{King}|\text{Face}) = \frac{P(\text{Face}|\text{King}).P(\text{King})}{P(\text{Face})}$$

Bayes' Theorem Example



$$P(\text{King}) = 4/52 = 1/13$$

$$P(\text{King}|\text{Face}) = \frac{P(\text{Face}|\text{King}).P(\text{King})}{P(\text{Face})}$$

$$P(\text{Face}|\text{King}) = 1$$

Bayes' Theorem Example



$$P(\text{King}) = 4/52 = 1/13$$

$$P(\text{King}|\text{Face}) = \frac{P(\text{Face}|\text{King}).P(\text{King})}{P(\text{Face})}$$

$$P(\text{Face}|\text{King}) = 1$$

$$P(\text{Face}) = 12/52 = 3/13$$

Bayes' Theorem Example



$$P(\text{King}) = 4/52 = 1/13$$

$$P(\text{King}|\text{Face}) = \frac{P(\text{Face}|\text{King}).P(\text{King})}{P(\text{Face})}$$

$$P(\text{Face}|\text{King}) = 1$$

$$= \frac{1.(1/13)}{3/13} = 1/3$$

$$P(\text{Face}) = 12/52 = 3/13$$

Bayes' Theorem Example

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

$$P(B|A) = \frac{P(B \cap A)}{P(A)}$$

$$P(A \cap B) = P(A|B) \cdot P(B) = P(B|A) \cdot P(A)$$

$$= P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$$

Bayes' Theorem Proof

Likelihood

How probable is the evidence
Given that our hypothesis is true?

$$P(H|E) = \frac{P(E|H) \cdot P(H)}{P(E)}$$

Posterior

How probable is our Hypothesis
Given the observed evidence?
(Not directly computable)

Prior

How probable was our hypothesis
Before observing the evidence?

Marginal

How probable is the new evidence
Under all possible hypothesis?

Naïve Bayes: Working

Classification Steps

Day	Outlook	Humidity	Wind	Play
D1	Sunny	High	Weak	No
D2	Sunny	High	Strong	No
D3	Overcast	High	Weak	Yes
D4	Rain	High	Weak	Yes
D5	Rain	Normal	Weak	Yes
D6	Rain	Normal	Strong	No
D7	Overcast	Normal	Strong	Yes
D8	Sunny	High	Weak	No
D9	Sunny	Normal	Weak	Yes
D10	Rain	Normal	Weak	Yes
D11	Sunny	Normal	Strong	Yes
D12	Overcast	High	Strong	Yes
D13	Overcast	Normal	Weak	Yes
D14	Rain	High	Strong	No



Outlook	Frequency Table		Play	
	Yes	No	Sunny	2
Overcast	4	0	Overcast	0
Rainy	3	2	Rainy	2

Humidity	Frequency Table		Play	
	Yes	No	High	4
Normal	6	1	Normal	1

Wind	Frequency Table		Play	
	Yes	No	Strong	2
Weak	3	3	Weak	3

Classification Steps

Likelihood Table		Play		
		Yes	No	
Outlook	Sunny	3/10	2/4	5/14
	Overcast	4/10	0/4	4/14
	Rainy	3/10	2/4	5/14
		10/14	4/14	

$P(x|c) = P(\text{Sunny}|\text{Yes}) = 3/10 = 0.3$

$P(x) = P(\text{Sunny}) = 5/14 = 0.36$

$P(c) = P(\text{Yes}) = 10/14 = 0.71$

Likelihood of 'Yes' given Sunny is

$$P(c|x) = P(\text{Yes}|\text{Sunny}) = P(\text{Sunny}|\text{Yes}) * P(\text{Yes}) / P(\text{Sunny}) = (0.3 \times 0.71) / 0.36 = 0.591$$

Similarly Likelihood of 'No' given Sunny is

$$P(c|x) = P(\text{No}|\text{Sunny}) = P(\text{Sunny}|\text{No}) * P(\text{No}) / P(\text{Sunny}) = (0.4 \times 0.36) / 0.36 = 0.40$$

Classification Steps

Likelihood table for Humidity

Likelihood Table		Play		
		Yes	No	
Humidity	High	3/9	4/5	7/14
	Normal	6/9	1/5	7/14
		9/14	5/14	

$$P(\text{Yes}|\text{High}) = 0.33 \times 0.6 / 0.5 = 0.42$$

$$P(\text{No}|\text{High}) = 0.8 \times 0.36 / 0.5 = 0.58$$

Likelihood table for Wind

Likelihood Table		Play		
		Yes	No	
Wind	Weak	6/9	2/5	8/14
	Strong	3/9	3/5	6/14
		9/14	5/14	

$$P(\text{Yes}|\text{Weak}) = 0.67 \times 0.64 / 0.57 = 0.75$$

$$P(\text{No}|\text{Weak}) = 0.4 \times 0.36 / 0.57 = 0.25$$

Classification Steps

Suppose we have a day with the following values

Outlook	=	Rain
Humidity	=	High
Wind	=	Weak
Play	=	?

$$\begin{aligned}\text{Likelihood of 'Yes' on that Day} &= P(\text{Outlook} = \text{Rain} | \text{Yes}) * P(\text{Humidity} = \text{High} | \text{Yes}) * P(\text{Wind} = \text{Weak} | \text{Yes}) * P(\text{Yes}) \\ &= 2/9 * 3/9 * 6/9 * 9/14 = 0.0199\end{aligned}$$

$$\begin{aligned}\text{Likelihood of 'No' on that Day} &= P(\text{Outlook} = \text{Rain} | \text{No}) * P(\text{Humidity} = \text{High} | \text{No}) * P(\text{Wind} = \text{Weak} | \text{No}) * P(\text{No}) \\ &= 2/5 * 4/5 * 2/5 * 5/14 = 0.0166\end{aligned}$$

Classification Steps

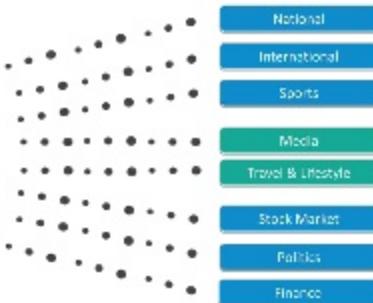
$$P(\text{Yes}) = 0.0199 / (0.0199 + 0.0166) = 0.55$$

$$P(\text{No}) = 0.0166 / (0.0199 + 0.0166) = 0.45$$

Our model predicts that
there is a 55% chance
there will be game
tomorrow



Industrial Use Cases



News Categorization

Weather Predictions

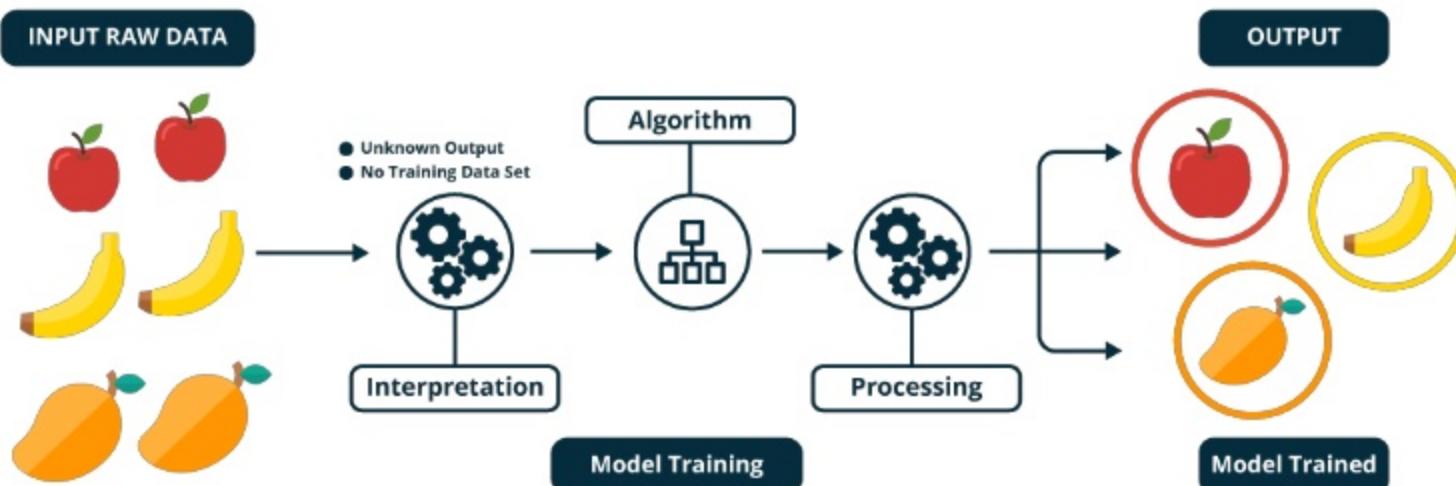


Spam Filtering

Demo

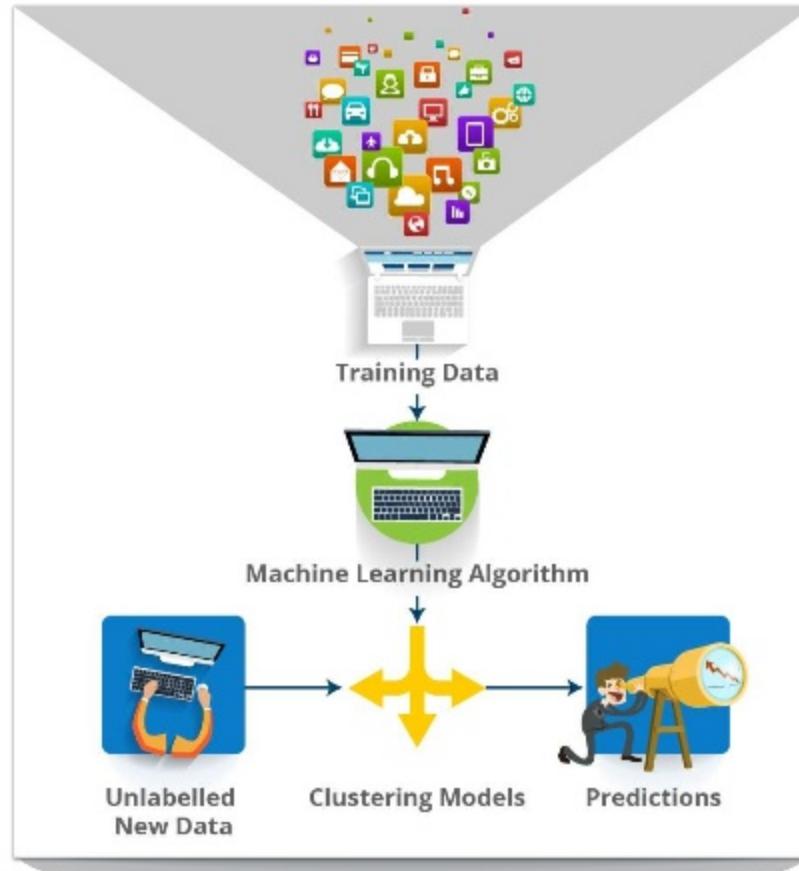
Unsupervised Learning

Unsupervised learning is a type of machine learning algorithm used to draw inferences from datasets consisting of input data **Without labelled responses**



Unsupervised Learning: Process Flow

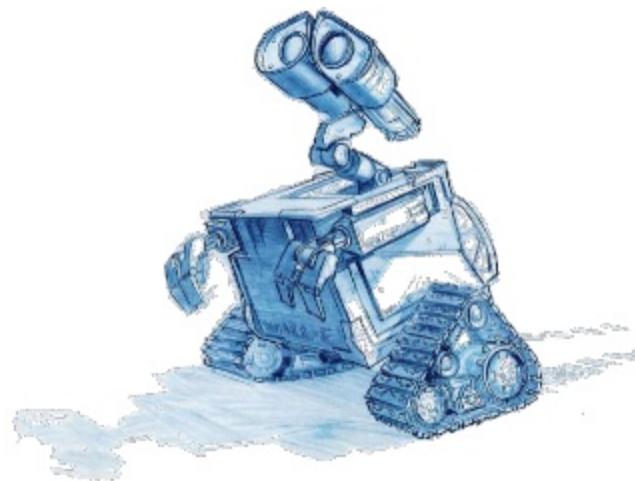
Training data is collection of information without any label



What is Clustering?

“Clustering is the process of dividing the datasets into groups, consisting of similar data-points”

It means grouping of objects based on the information found in the data, describing the objects or their relationship



Why is Clustering Used?

The goal of clustering is to determine the intrinsic grouping
in a set of **Unlabelled Data**



Sometimes, Partitioning is the goal

Where is it used?



Retail Store



Banking



Insurance
Companies



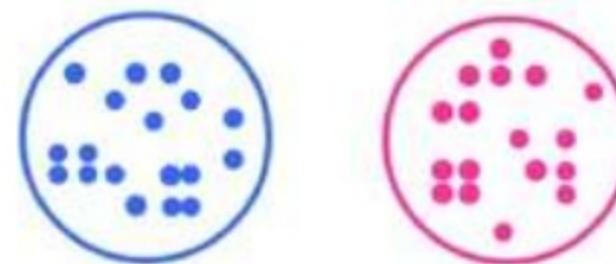
NETFLIX

Recommended Movies



Types of Clustering

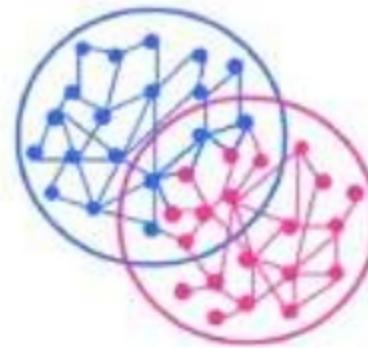
- Exclusive Clustering
- Overlapping Clustering
- Hierarchical Clustering



K-Means Clustering

Types of Clustering

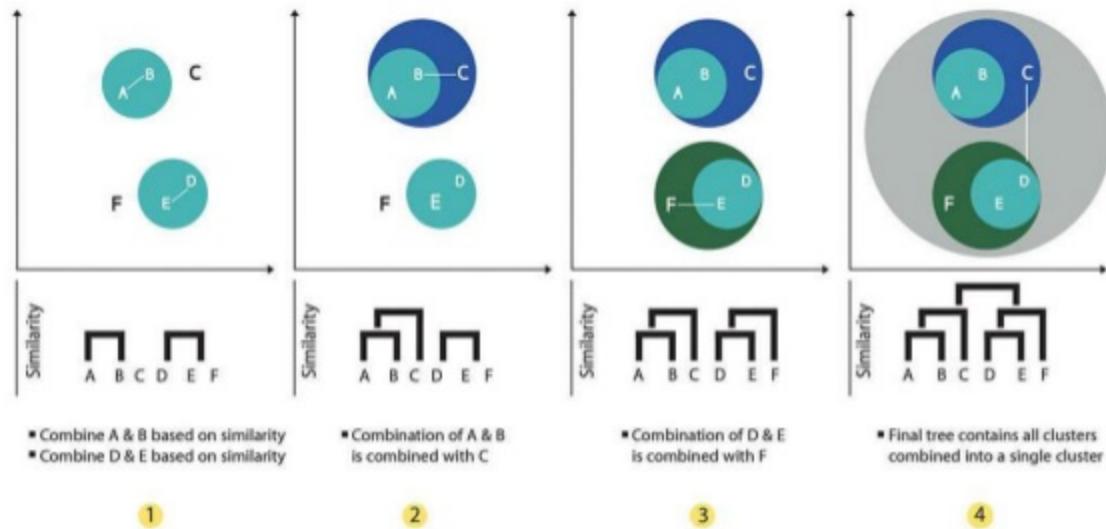
- Exclusive Clustering
- Overlapping Clustering
- Hierarchical Clustering



C-Means Clustering

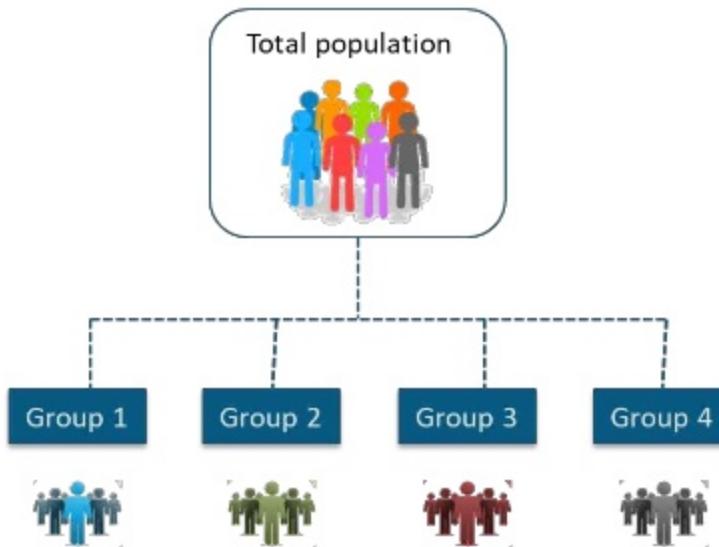
Types of Clustering

- Exclusive Clustering
- Overlapping Clustering
- Hierarchical Clustering

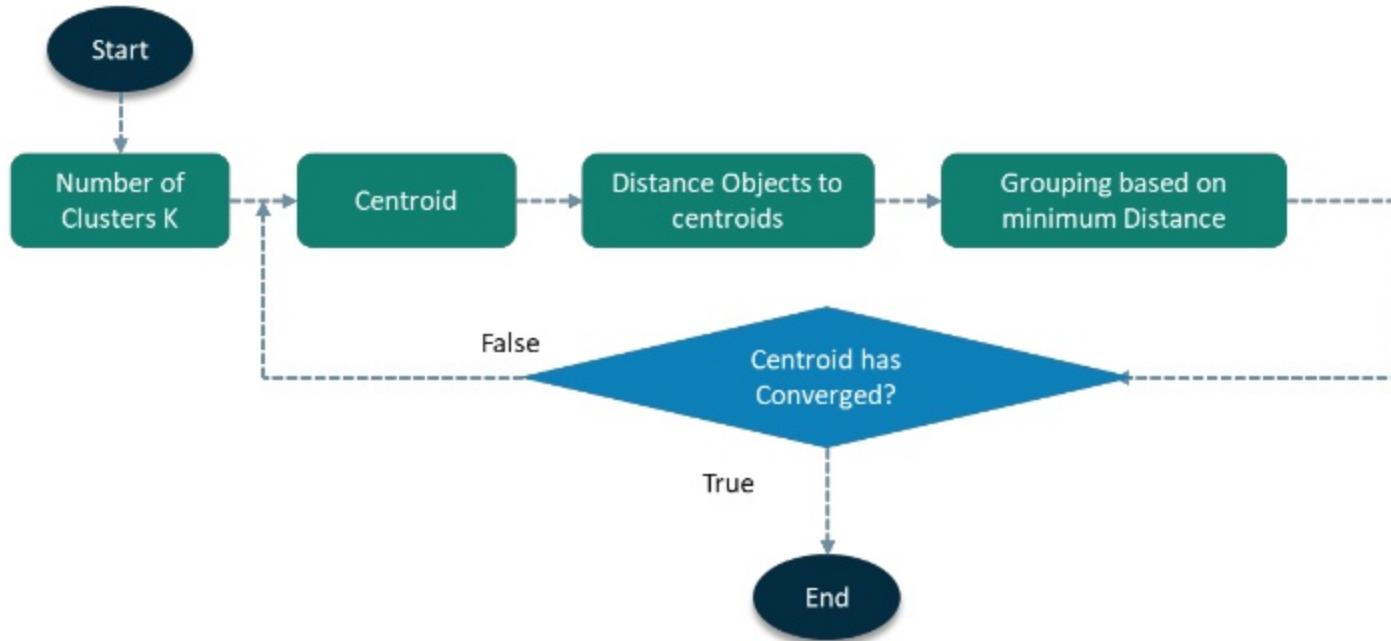


K-Means Clustering

The process by which objects are classified into a **predefined number of groups** so that they are as much dissimilar as possible from one group to another group, but as much **similar as possible** within each group.



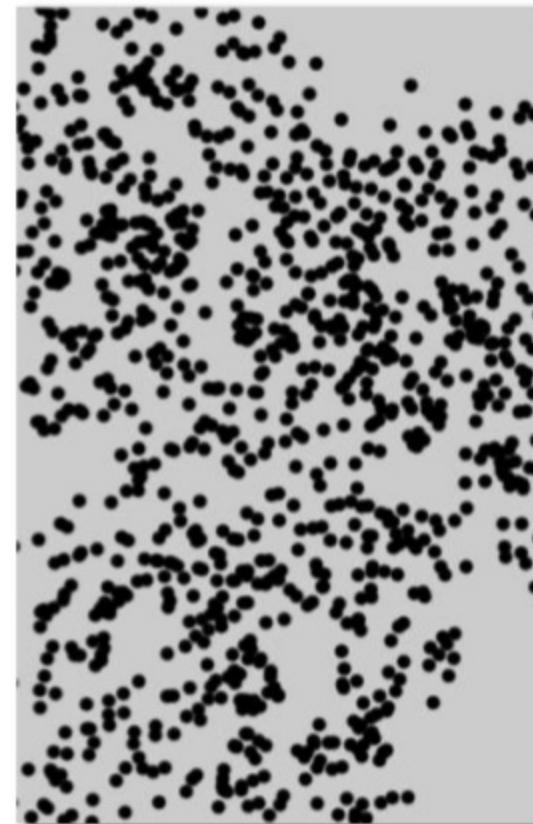
K-Means Algorithm Working



K-Means Clustering : Steps

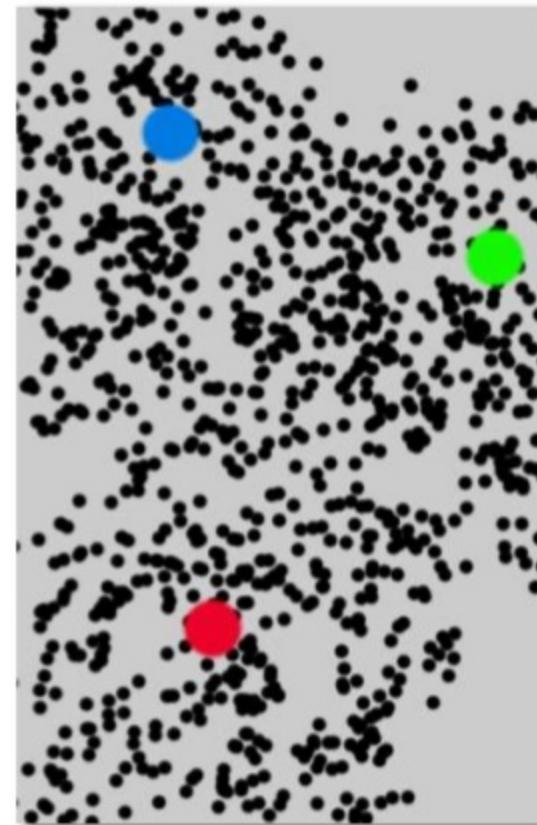
- 1 First we need to decide the number of clusters to be made. (Guessing)

Let's assume , Number of clusters = 3



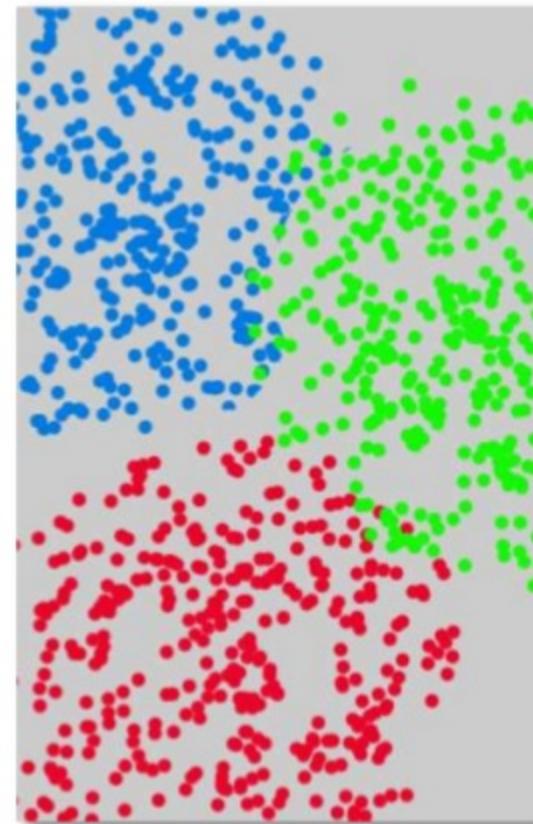
K-Means Clustering : Steps

- ① First we need to decide the number of clusters to be made. (Guessing)
- ② Then we provide centroids of all the clusters. (Guessing)



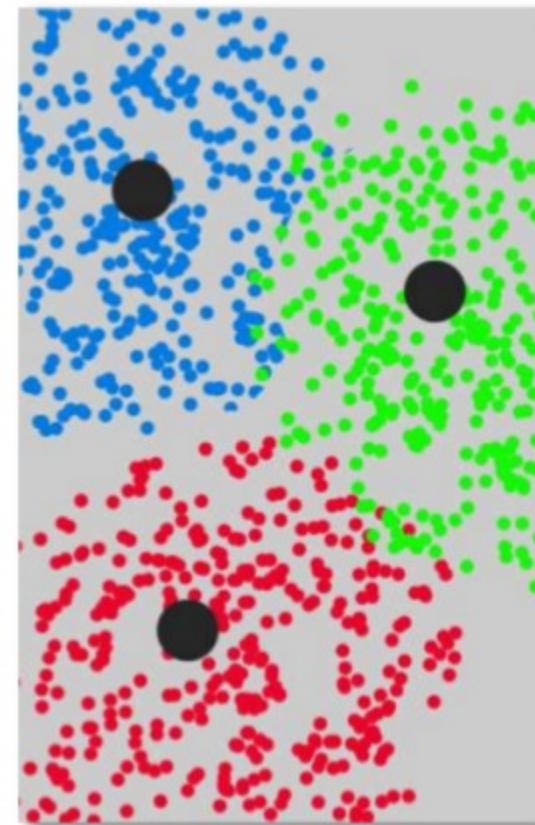
K-Means Clustering : Steps

- ① First we need to decide the number of clusters to be made. (Guessing)
- ② Then we provide centroids of all the clusters. (Guessing)
- ③ The Algorithm calculates Euclidian distance of the points from each centroid and assigns the point to the closest cluster.



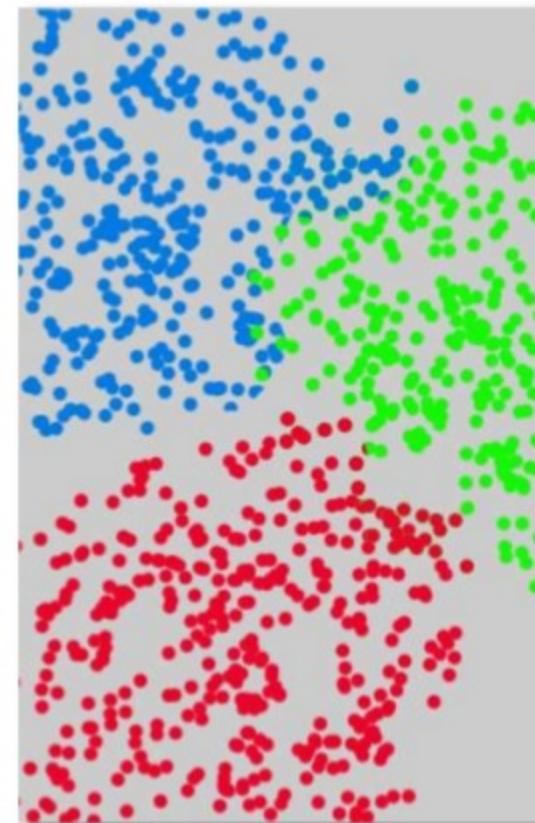
K-Means Clustering : Steps

- 1 First we need to decide the number of clusters to be made. (Guessing)
- 2 Then we provide centroids of all the clusters. (Guessing)
- 3 The Algorithm calculates Euclidian distance of the points from each centroid and assigns the point to the closest cluster.
- 4 Next the Centroids are calculated again, when we have our new cluster.



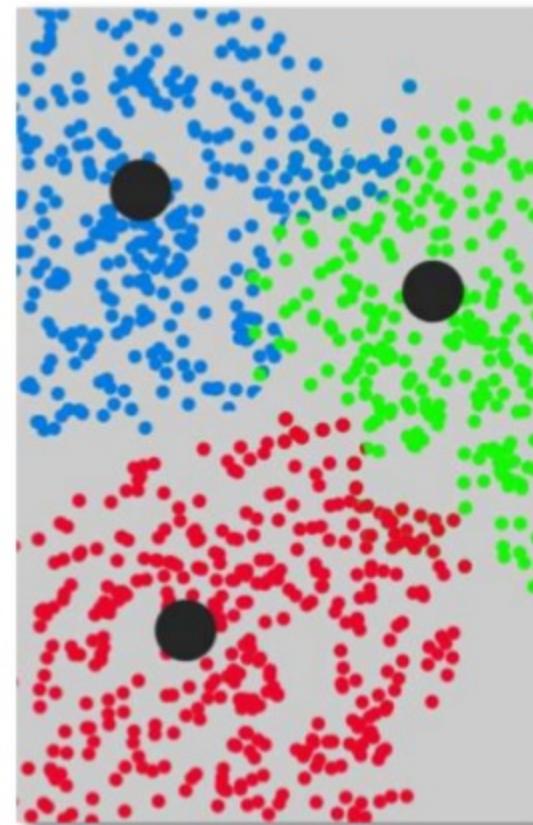
K-Means Clustering : Steps

- 1 First we need to decide the number of clusters to be made. (Guessing)
- 2 Then we provide centroids of all the clusters. (Guessing)
- 3 The Algorithm calculates Euclidian distance of the points from each centroid and assigns the point to the closest cluster.
- 4 Next the Centroids are calculated again, when we have our new cluster.
- 5 The distance of the points from the centre of clusters are calculated again and points are assigned to the closest cluster.



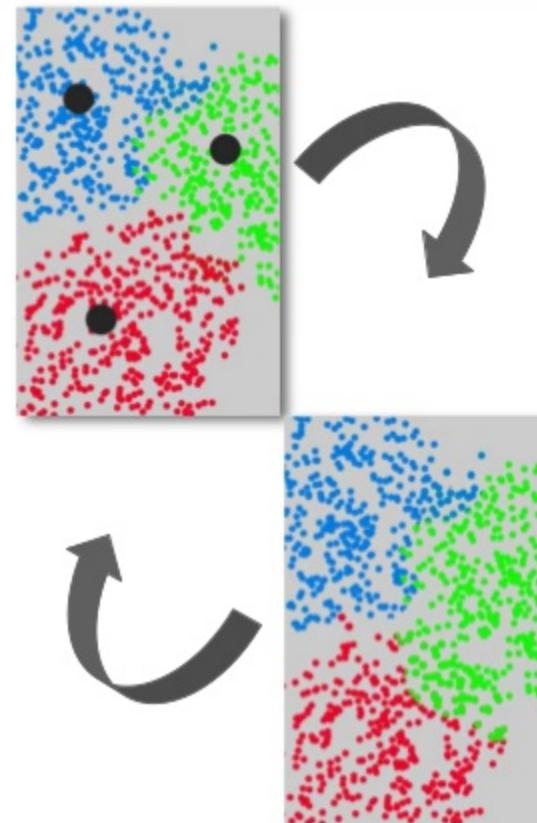
K-Means Clustering : Steps

- 1 First we need to decide the number of clusters to be made. (Guessing)
- 2 Then we provide centroids of all the clusters. (Guessing)
- 3 The Algorithm calculates Euclidian distance of the points from each centroid and assigns the point to the closest cluster.
- 4 Next the Centroids are calculated again, when we have our new cluster.
- 5 The distance of the points from the centre of clusters are calculated again and points are assigned to the closest cluster.
- 6 And then again the new centroid for the cluster is calculated.



K-Means Clustering : Steps

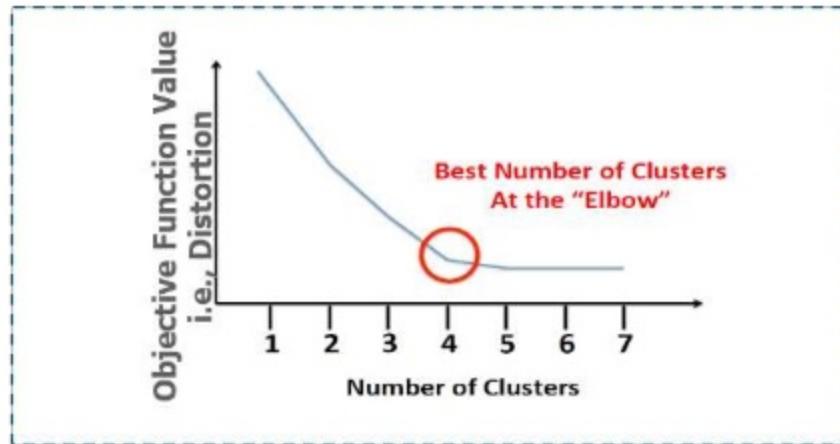
- 1 First we need to decide the number of clusters to be made. (Guessing)
- 2 Then we provide centroids of all the clusters. (Guessing)
- 3 The Algorithm calculates Euclidian distance of the points from each centroid and assigns the point to the closest cluster.
- 4 Next the Centroids are calculated again, when we have our new cluster.
- 5 The distance of the points from the centre of clusters are calculated again and points are assigned to the closest cluster.
- 6 And then again the new centroid for the cluster is calculated.
- 7 These steps are repeated until we have a repetition in centroids or new centroids are very close to the previous ones.



How to Decide the number of Clusters

The Elbow Method :

First of all, compute the sum of squared error (SSE) for some values of k (for example 2, 4, 6, 8, etc.). The SSE is defined as the sum of the squared distance between each member of the cluster and its centroid. Mathematically:



$$SSE = \sum_{i=1}^K \sum_{x \in c_i} dist(x, c_i)^2$$

Pros and Cons: K-Means Clustering



- Simple, understandable
- Items automatically assigned to clusters

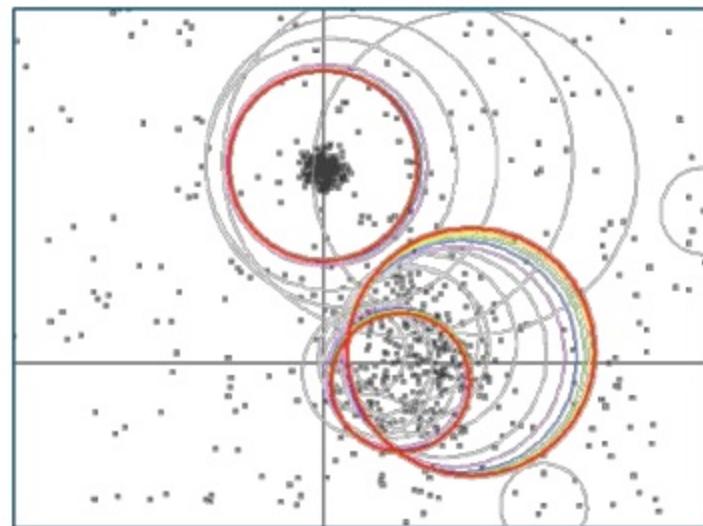
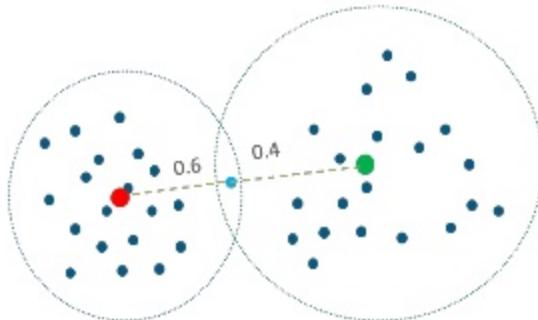
- Must define number of clusters
- All items forced into clusters
- Unable to handle noisy data and outliers



Fuzzy C – Means Clustering

Fuzzy C-Means is an extension of K-Means, the popular simple clustering technique

Fuzzy clustering (also referred to as soft clustering) is a form of Clustering in which each data point can belong to more than one cluster



Pros and Cons: C-Means Clustering



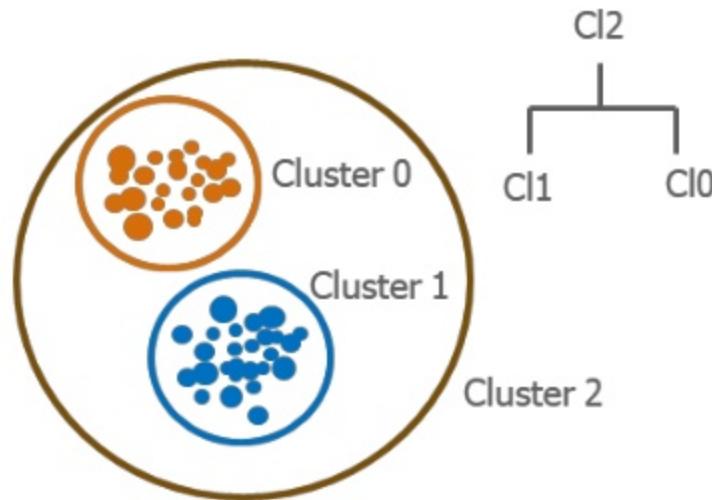
- Allows a data point to be in multiple clusters
- A more natural representation of the behaviour of genes
- Genes usually are involved in multiple functions

- Need to define c , the number of clusters
- Need to determine membership cut-off value
- Clusters are sensitive to initial assignment of centroids
- Fuzzy c-means is not a deterministic algorithm



Hierarchical Clustering

Hierarchical clustering is an alternative approach which builds a hierarchy from the bottom-up, and doesn't require us to specify the number of clusters beforehand



Pros and Cons: Hierarchical Clustering



- No assumption of a particular number of clusters
- May corresponds to meaningful taxonomies

- Once a decision is made to combine two clusters, it can't be undone
- Too slow for large datasets



Demo

Market Basket Analysis

Market basket analysis explains the combinations of products that frequently co-occur in transactions.



Market Basket Analysis algorithms

1. Association Rule Mining
2. Apriori

Association Rule Mining

Association rule mining is a technique that shows how items are associated to each other.

Example:

Customer who purchase bread have a 60% likelihood of also purchasing jam.



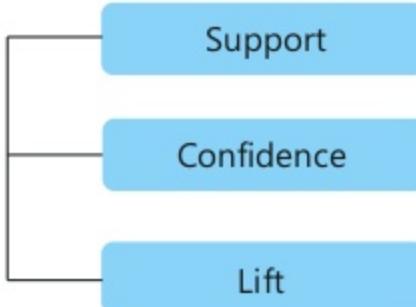
Customer who purchase laptops are more likely to purchase laptop bags.

Association Rule Mining

Example of Association rule

A \Rightarrow B

- It means that if a person buys item A then he will also buy item B
- Three common ways to measure association:



Association Rule Mining

Support gives fraction of transactions which contains the item A and B

$$Support = \frac{freq(A, B)}{N}$$

Confidence gives how often the items A & B occur together, given no. of times A occurs

$$Confidence = \frac{freq(A, B)}{freq(A)}$$

Lift indicates the strength of a rule over the random co-occurrence of A and B

$$Lift = \frac{Support}{Supp(A) \times Supp(B)}$$

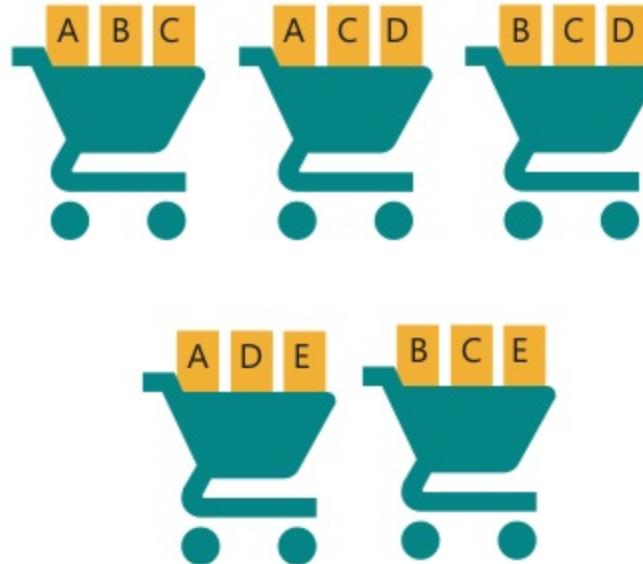
Association Rule Mining Example

Set of items {A, B, C, D, E}

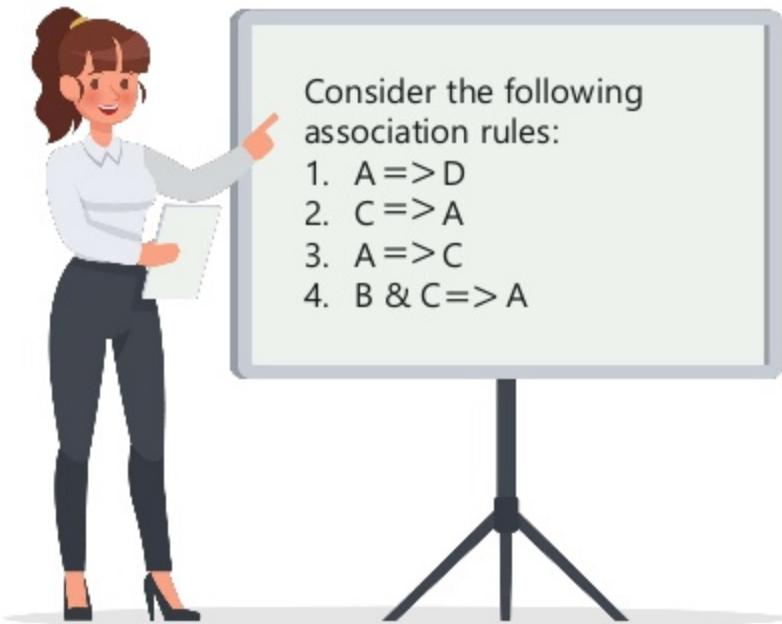
Set of transactions {T1, T2, T3, T4, T5}

Transactions at a local store

T1	A	B	C
T2	A	C	D
T3	B	C	D
T4	A	D	E
T5	B	C	E



Association Rule Mining Example



Calculate support, confidence and lift for these rules:

Rule	Support	Confidence	Lift
$A \Rightarrow D$	2/5	2/3	10/9
$C \Rightarrow A$	2/5	2/4	5/6
$A \Rightarrow C$	2/5	2/3	5/6
$B, C \Rightarrow A$	1/5	1/3	5/9

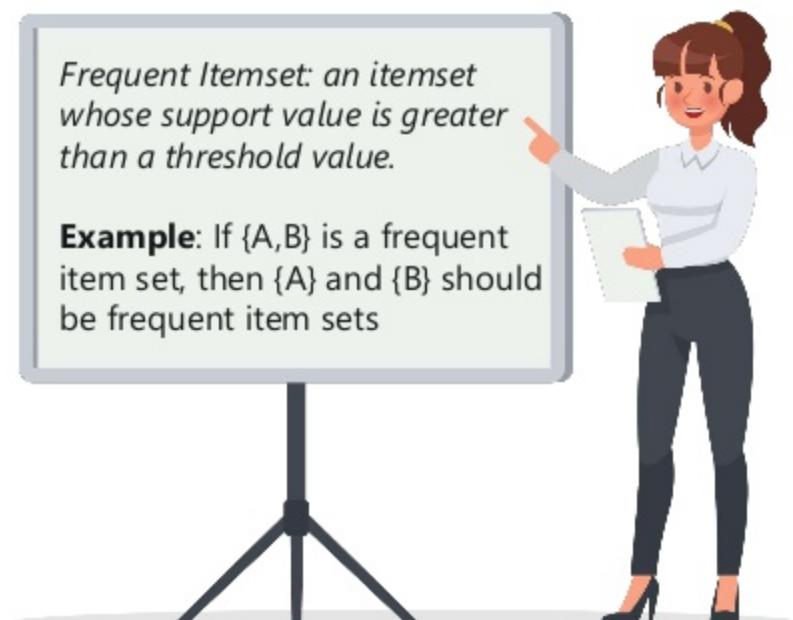
Apriori Algorithm

Apriori algorithm uses frequent item sets to generate association rules. It is based on the concept that a subset of a frequent itemset must also be a frequent itemset.



Frequent Itemset: an itemset whose support value is greater than a threshold value.

Example: If $\{A, B\}$ is a frequent item set, then $\{A\}$ and $\{B\}$ should be frequent item sets



Apriori Algorithm

Consider the following transactions:

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5

Now the first step is to build a list of item sets of size one by using this transactional dataset.



Note:

Min. support count = 2

Apriori Algorithm – First Iteration

Step 1: Create item sets of size one & calculate their support values

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5



Table: C1|

Itemset	Support
{1}	3
{2}	3
{3}	4
{4}	1
{5}	4



Table: F1|

Itemset	Support
{1}	3
{2}	3
{3}	4
{5}	4

Item sets with support value less than min. support value (i.e. 2) are eliminated

Apriori Algorithm – Second Iteration

Step 2: Create item sets of size two & calculate their support values.
All the combinations of item sets in F1| are used in this iteration

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5



Table: C2|

Itemset	Support
{1,2}	1
{1,3}	3
{1,5}	2
{2,3}	2
{2,5}	3
{3,5}	3



Table: F2|

Itemset	Support
{1,3}	3
{1,5}	2
{2,3}	2
{2,5}	3
{3,5}	3

Item sets with support less than 2 it are eliminated

Apriori Algorithm – Third Iteration

Step 3: Create item sets of size three & calculate their support values.
All the combinations of item sets in $F_2 \cup F_3$ are used in this iteration

Before calculating support values, let's perform pruning on the dataset!

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5



Table: C3	
Itemset	Support
{1,2,3}	
{1,2,5}	
{1,3,5}	
{2,3,5}	



Apriori Algorithm – Pruning

After the combinations are made, divide C3| item sets to check if there are any other subsets whose support is less than min support value.

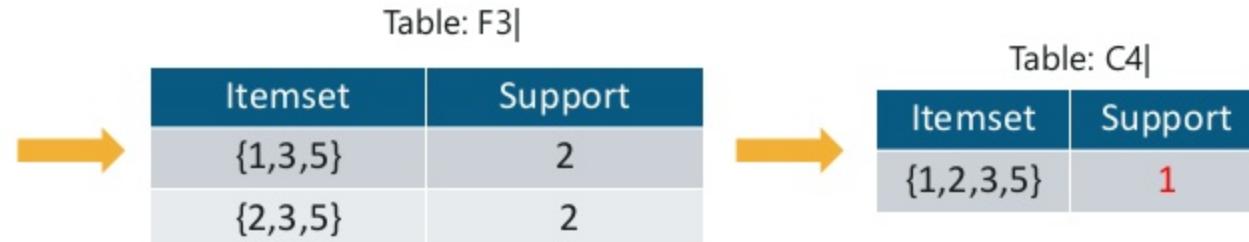
		Table: C3		Table: F2	
TID	Items	Itemset	In F2 ?	Itemset	Support
T1	1 3 4	{1,2,3}	NO	{1,3}	3
T2	2 3 5	{1,2},{1,3},{2,3}		{1,5}	2
T3	1 2 3 5	{1,2,5}	NO	{2,3}	2
T4	2 5	{1,2},{1,5},{2,5}		{2,5}	3
T5	1 3 5	{1,3,5}	YES	{3,5}	3
		{1,5},{1,3},{3,5}			
		{2,3,5}	YES		
		{2,3},{2,5},{3,5}			

If any of the subsets of these item sets are not there in F12 then we remove that itemset

Apriori Algorithm – Fourth Iteration

Using the item sets of C3|, create new itemset C4|.

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5



Since support of C4| is less than 2, stop and return to the previous itemset, i.e. C13

Apriori Algorithm – Subset Creation

Frequent Item set F3|

Itemset	Support
{1,3,5}	2
{2,3,5}	2

Let's assume our
minimum confidence
value is 60%



- Generate all non empty subsets for each frequent item set
 - ❖ For I = {1,3,5}, subsets are {1,3}, {1,5}, {3,5}, {1}, {3}, {5}
 - ❖ For I = {2,3,5}, subsets are {2,3}, {2,5}, {3,5}, {2}, {3}, {5}
- For every subsets S of I, output the rule:

S → (I-S) (S recommends I-S)

if $support(I)/support(S) \geq min_conf\ value$

Apriori Algorithm – Applying Rules

Applying rules to item sets of F3|:

1. {1,3,5}

- ✓ Rule 1: $\{1,3\} \rightarrow (\{1,3,5\} - \{1,3\})$ means $1 \& 3 \rightarrow 5$

Confidence = $\text{support}(1,3,5)/\text{support}(1,3) = 2/3 = 66.66\% > 60\%$

Rule 1 is selected

- ✓ Rule 2: $\{1,5\} \rightarrow (\{1,3,5\} - \{1,5\})$ means $1 \& 5 \rightarrow 3$

Confidence = $\text{support}(1,3,5)/\text{support}(1,5) = 2/2 = 100\% > 60\%$

Rule 2 is selected

- ✓ Rule 3: $\{3,5\} \rightarrow (\{1,3,5\} - \{3,5\})$ means $3 \& 5 \rightarrow 1$

Confidence = $\text{support}(1,3,5)/\text{support}(3,5) = 2/3 = 66.66\% > 60\%$

Rule 3 is selected

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5

Apriori Algorithm – Applying Rules

Applying rules to item sets of F3|:

1. {1,3,5}

- ✓ Rule 4: $\{1\} \rightarrow (\{1,3,5\} - \{1\})$ means $1 \rightarrow 3 \& 5$

Confidence = $\text{support}(1,3,5)/\text{support}(1) = 2/3 = 66.66\% > 60\%$

Rule 4 is selected

- ✓ Rule 5: $\{3\} \rightarrow (\{1,3,5\} - \{3\})$ means $3 \rightarrow 1 \& 5$

Confidence = $\text{support}(1,3,5)/\text{support}(3) = 2/4 = 50\% < 60\%$

Rule 5 is rejected

- ✓ Rule 6: $\{5\} \rightarrow (\{1,3,5\} - \{5\})$ means $5 \rightarrow 1 \& 3$

Confidence = $\text{support}(1,3,5)/\text{support}(3) = 2/4 = 50\% < 60\%$

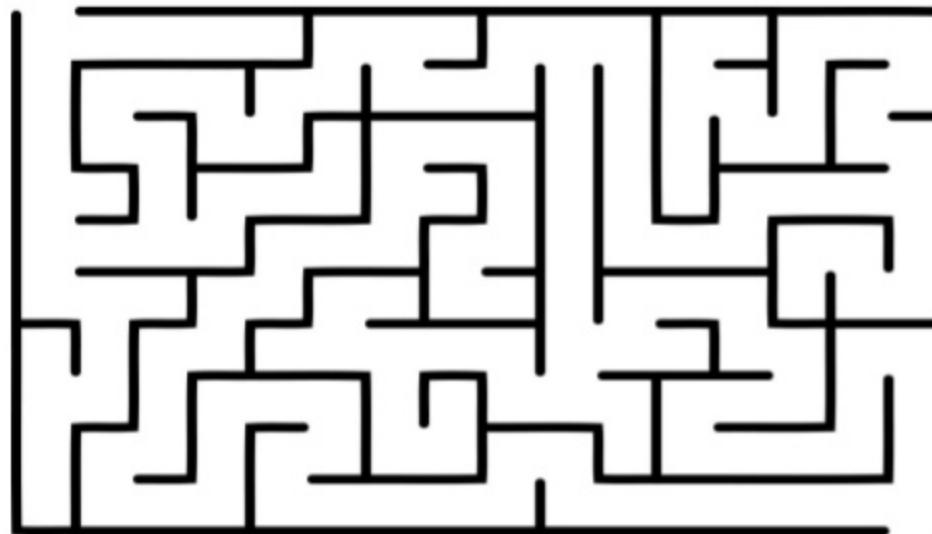
Rule 6 is rejected

TID	Items
T1	1 3 4
T2	2 3 5
T3	1 2 3 5
T4	2 5
T5	1 3 5

Demo

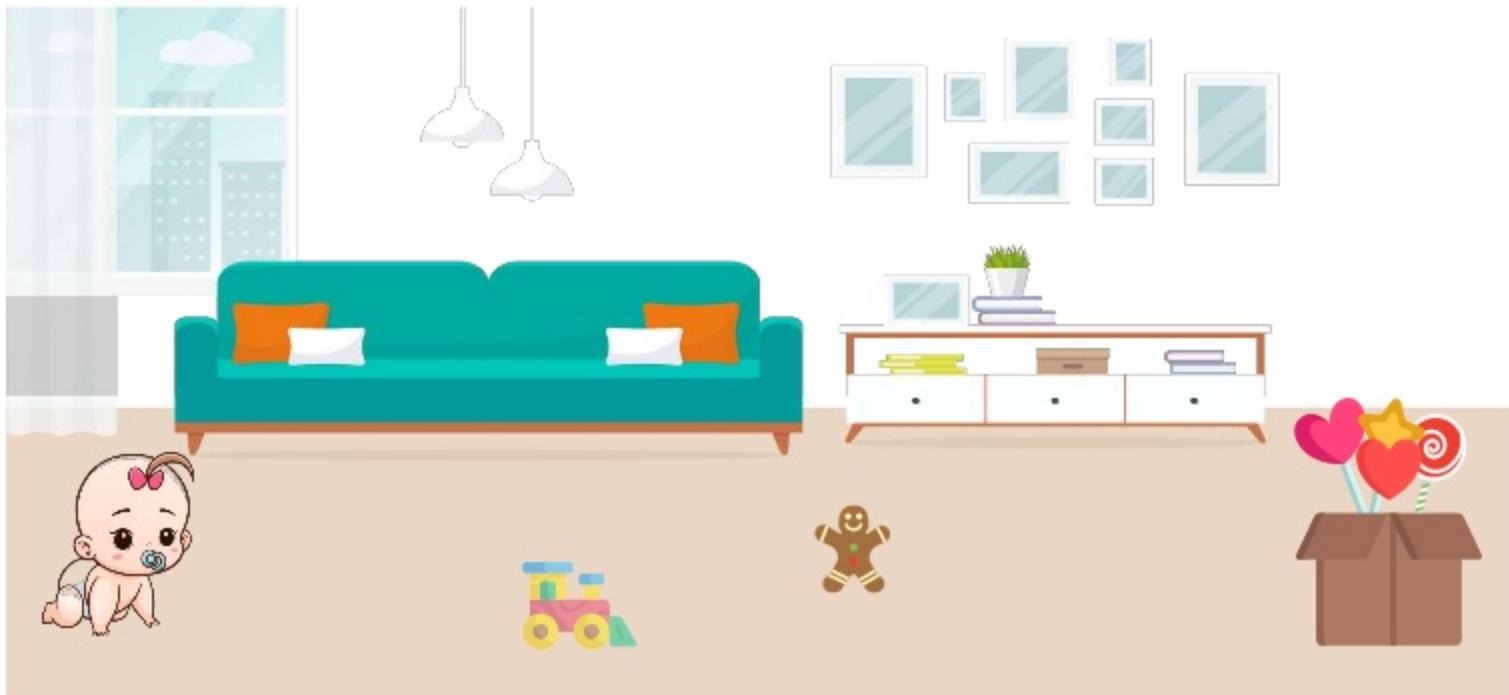
What is Reinforcement Learning?

Reinforcement learning is a type of Machine Learning where an agent learns to behave in an environment by performing actions and seeing the results



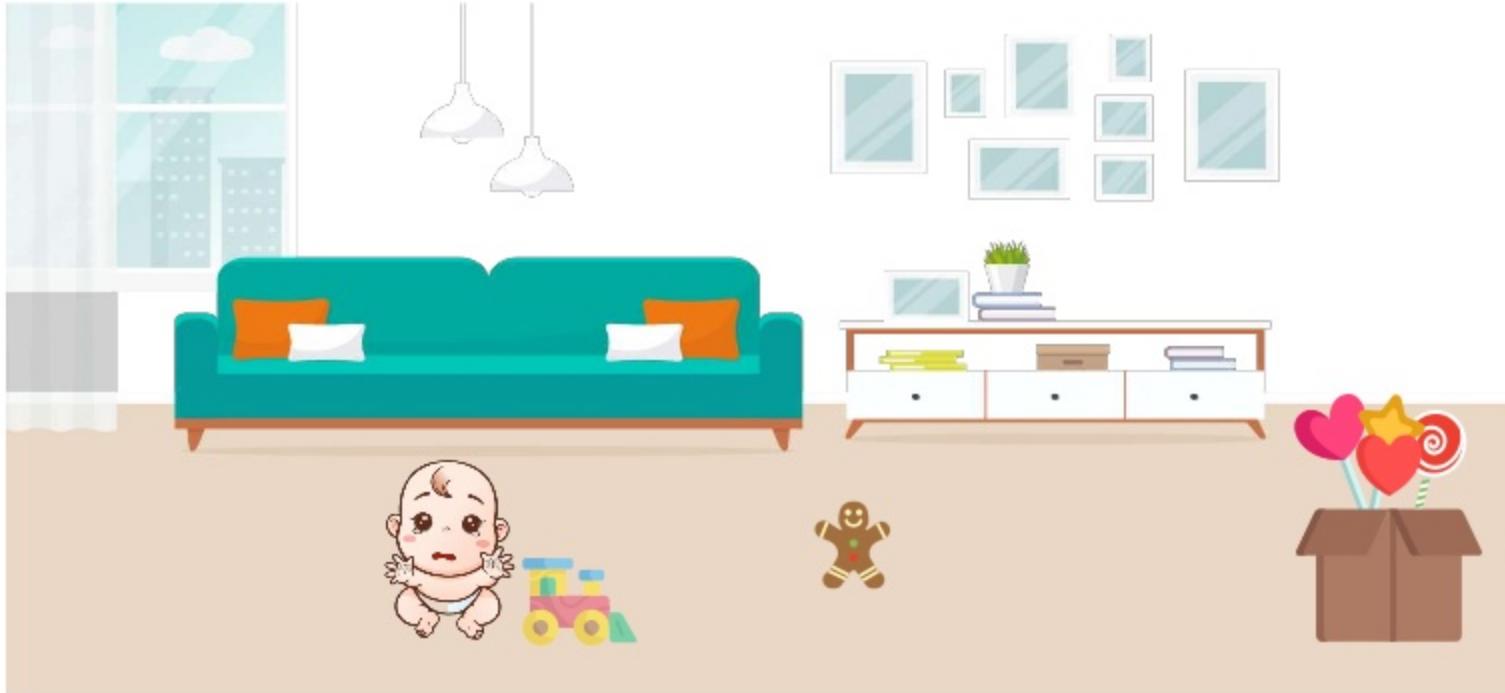
Analogy

Scenario 1: Baby starts crawling and makes it to the candy



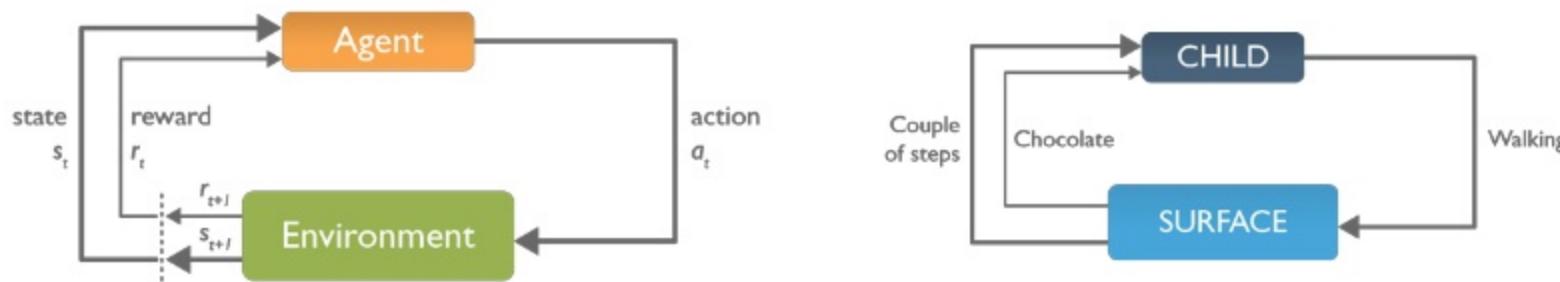
Analogy

Scenario 2: Baby starts crawling but falls due to some hurdle in between



Analogy

- The child is an **agent** trying to manipulate the **environment** (which is the surface on which it walks)
- Taking actions (viz walking) and he/she tries to go from one **state** (viz each step he/she takes) to another
- The child gets a **reward** (let's say chocolate) when he/she accomplishes a **submodule of the task** (viz taking couple of steps)



- Will not receive any chocolate (**negative reward**) when he/she is not able to walk

Reinforcement Learning Definitions

Model of the Environment
Used for planning & if know the current state and action then predict the resultant next state and next reward

Value Function

Value of a state is the total amount of reward an agent can expect to accumulate over the future

Agent

Intelligent programs



Reward Function

Could be +1 or any other value, indicating, what's good in an immediate sense

Environment

An external condition

Policy

A mapping from state to actions defining agent's behavior at a particular time

Reinforcement Learning Definitions



Agent: The RL algorithm that learns from trial and error



Action (A): All the possible steps that the agent can take



State (S): Current condition returned by the environment

Reinforcement Learning Definitions



Reward (R): An instant return from the environment to appraise the last action



Policy (π): The approach that the agent uses to determine the next action based on the current state



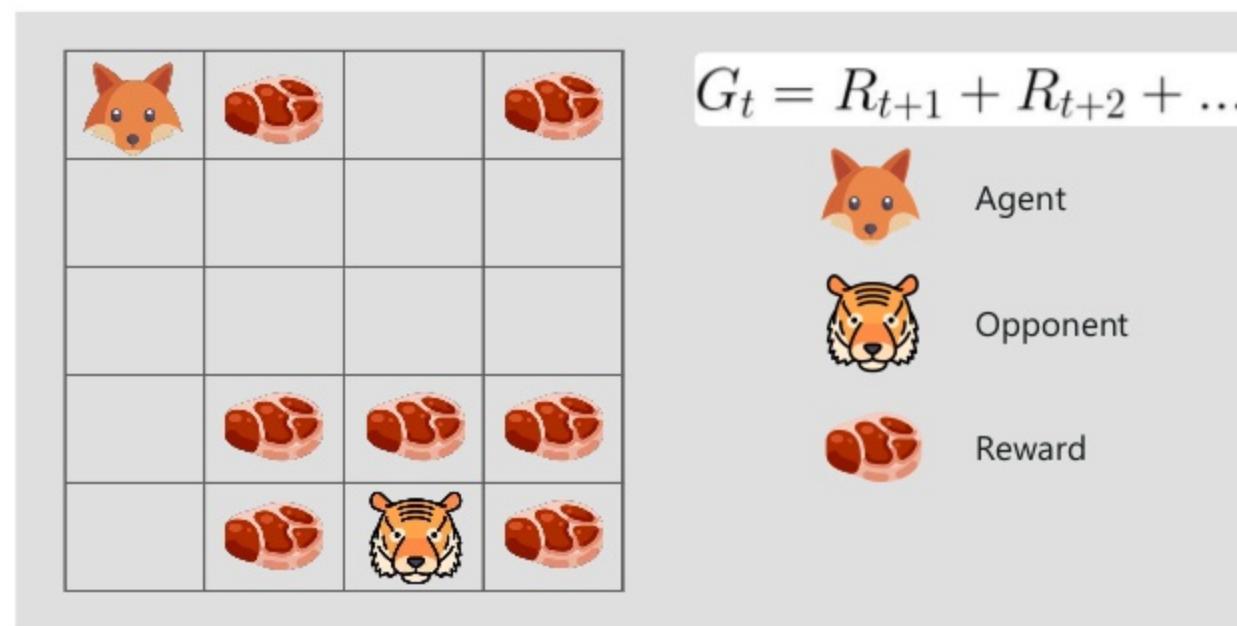
Value (V): The expected long-term return with discount, as opposed to the short-term reward R



Action-value (Q): This is similar to Value, except, it takes an extra parameter, the current action (A)

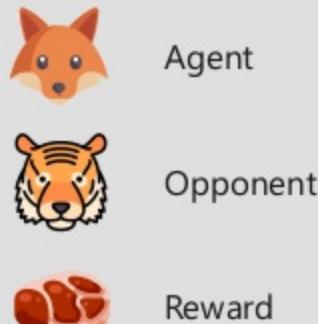
Reward Maximization

Reward maximization theory states that, a RL agent must be trained in such a way that, he takes the best action so that the reward is maximum.



Exploration & Exploitation

Exploitation is about using the already known exploited information to heighten the rewards
Exploration is about exploring and capturing more information about an environment



$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \text{ where } \gamma \in [0, 1)$$

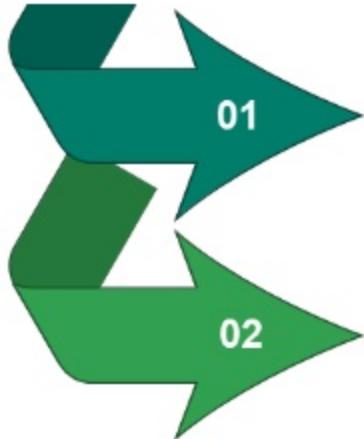
$$R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots$$

The K-Armed Bandit Problem



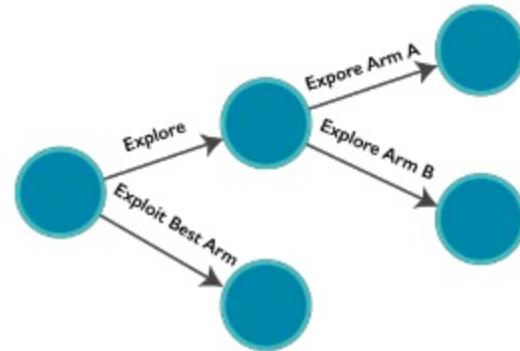
- The K-armed bandit is a metaphor representing a casino slot machine with k pull levers (or arms)
- The user or customer pulls any one of the levers to win a predefined reward
- The objective is to select the lever that will provide the user with the highest reward

The Epsilon Greedy Algorithm



Tries to be fair to the two opposite goals of exploration and exploitation by using a mechanism of flipping a coin

Just like, if you flip a coin and it comes up head you should explore for a moment but if it comes up tails, you should exploit



Takes whatever action seems best at the present moment

The Epsilon Greedy Algorithm



With probability $1 - \text{epsilon}$, the Epsilon-Greedy algorithm exploits the best known option



With probability $\text{epsilon} / 2$, the Epsilon-Greedy algorithm explores the best known option



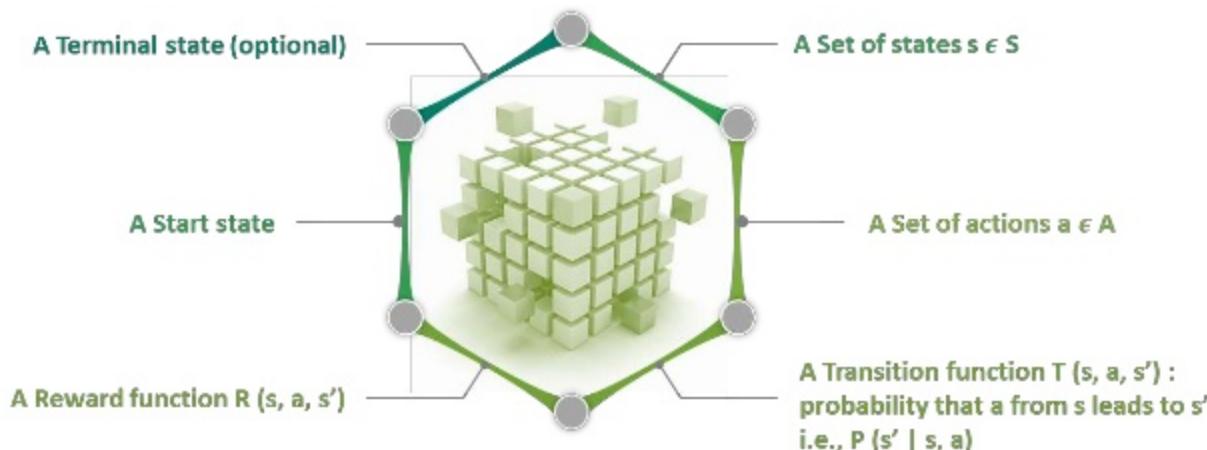
With probability $\text{epsilon} / 2$, the Epsilon-Greedy algorithm explores the worst known option



Markov Decision Process

The mathematical approach for mapping a solution in reinforcement learning is called *Markov Decision Process (MDP)*

The following parameters are used to attain a solution:

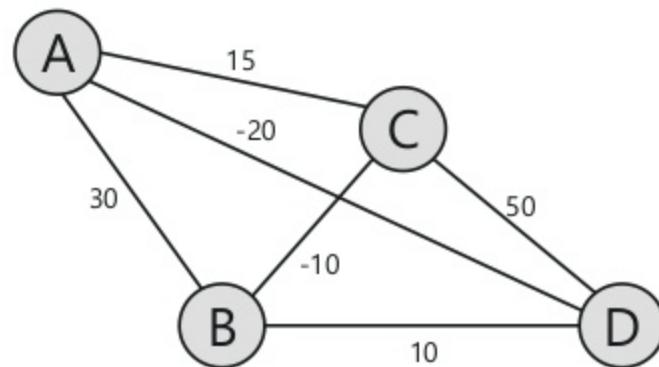


Markov Decision Process- Shortest Path Problem

Goal: Find the shortest path between A and D with minimum possible cost

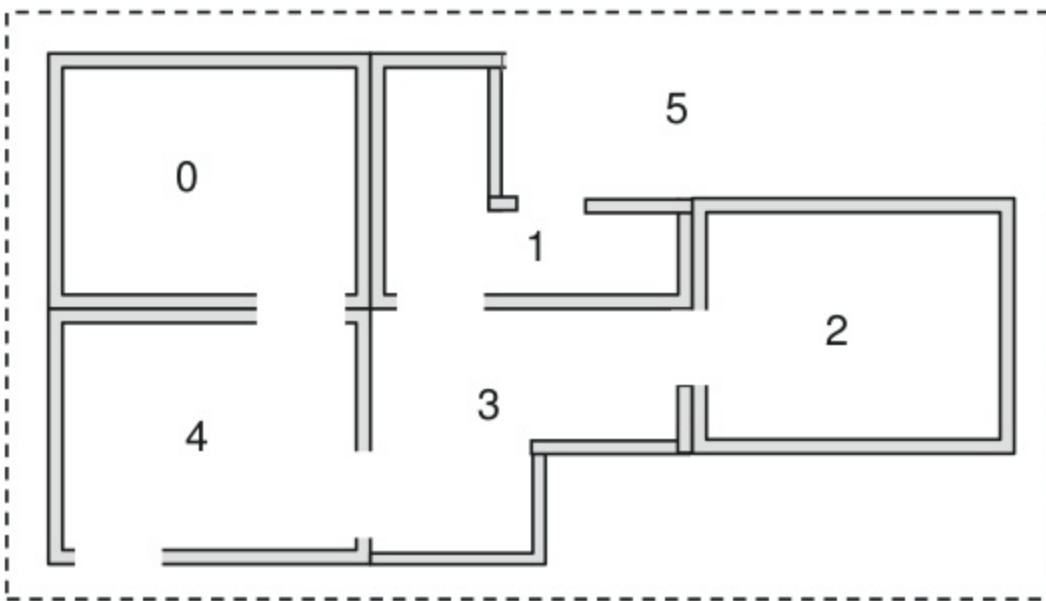
In this problem,

- Set of states are denoted by nodes i.e. {A, B, C, D}
- Action is to traverse from one node to another {A -> B, C -> D}
- Reward is the cost represented by each edge
- Policy is the path taken to reach the destination {A -> C -> D}



Understanding Q-Learning With An Example

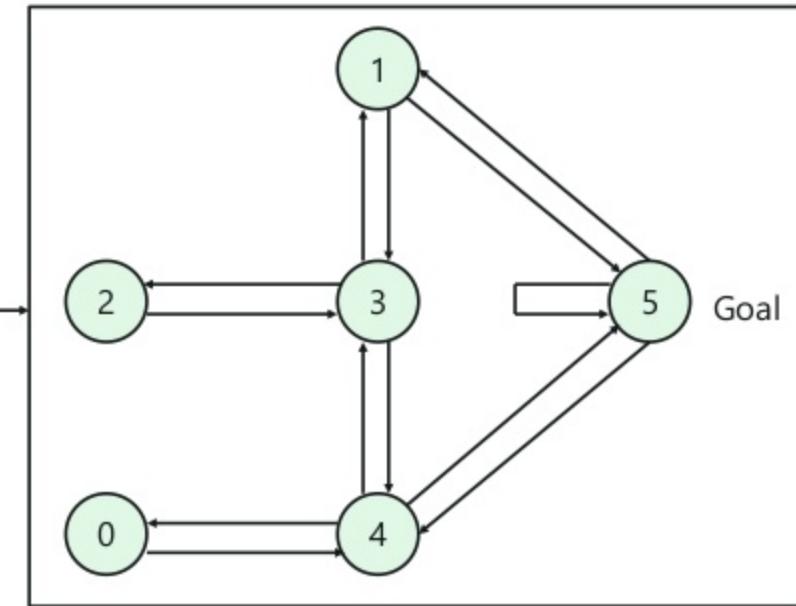
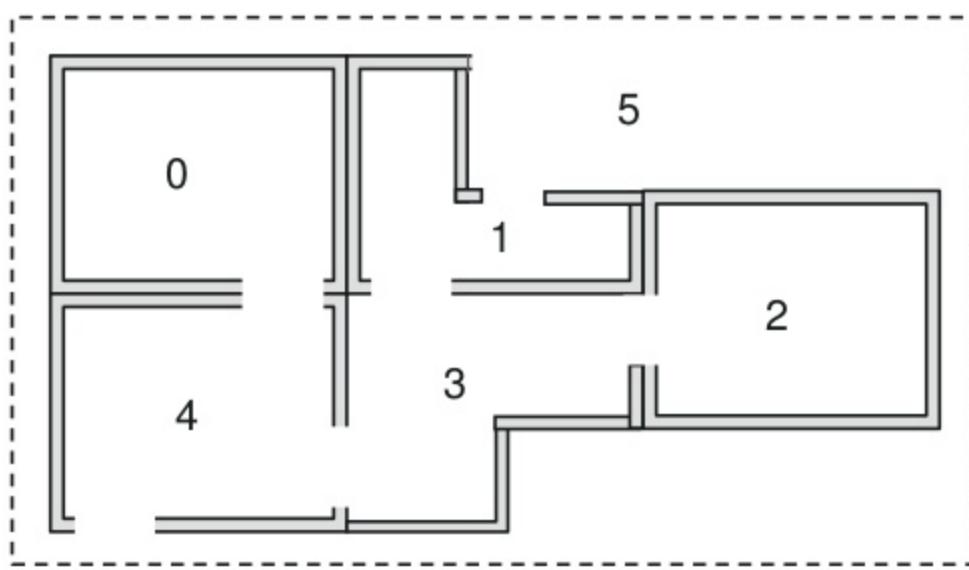
Place an agent in any one of the rooms (0,1,2,3,4) and the goal is to reach outside the building (room 5)



- 5 rooms in a building connected by doors
- each room is numbered 0 through 4
- The outside of the building can be thought of as one big room (5)
- Doors 1 and 4 lead into the building from room 5 (outside)

Understanding Q-Learning With An Example

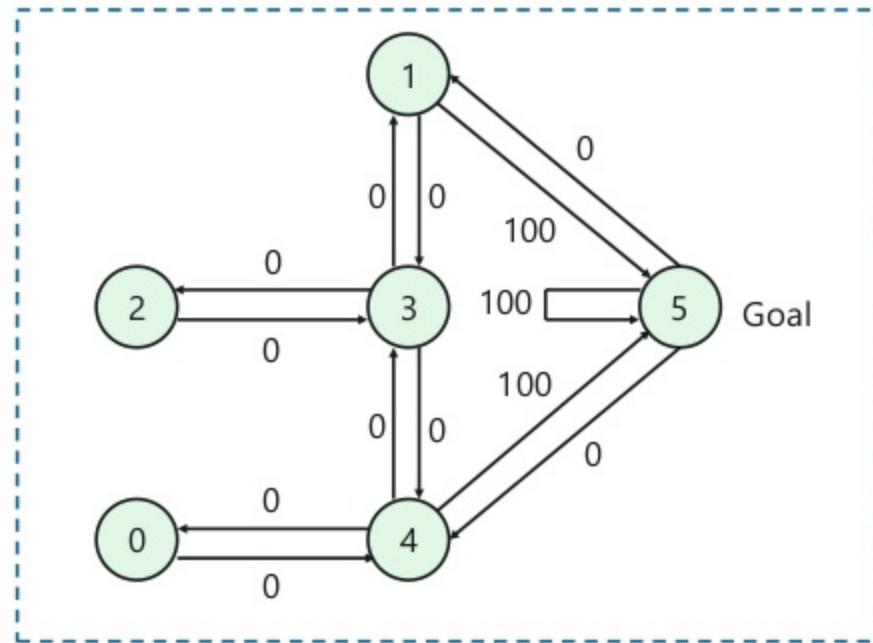
Let's represent the rooms on a graph, each room as a node, and each door as a link



Understanding Q-Learning With An Example

Next step is to associate a reward value to each door:

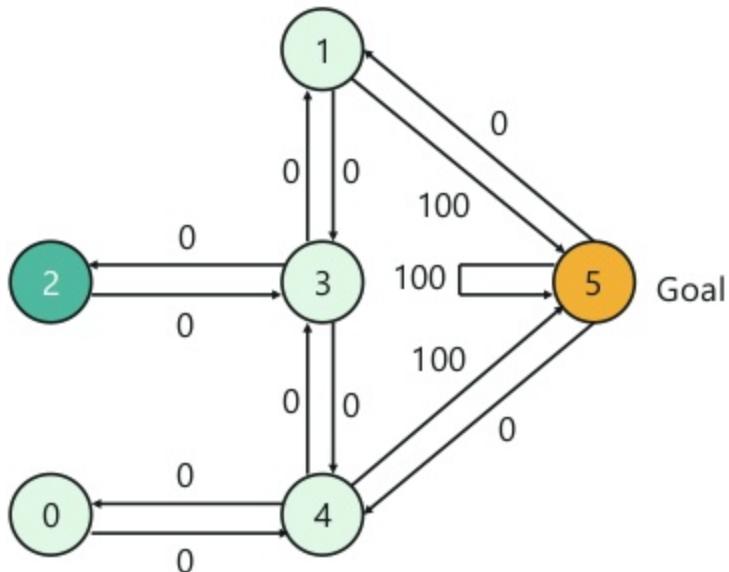
- doors that lead directly to the goal have a reward of 100
- Doors not directly connected to the target room have zero reward
- Because doors are two-way, two arrows are assigned to each room
- Each arrow contains an instant reward value



Understanding Q-Learning With An Example

The terminology in Q-Learning includes the terms state and action:

- Room (including room 5) represents a state
- agent's movement from one room to another represents an action
- In the figure, a state is depicted as a node, while "action" is represented by the arrows

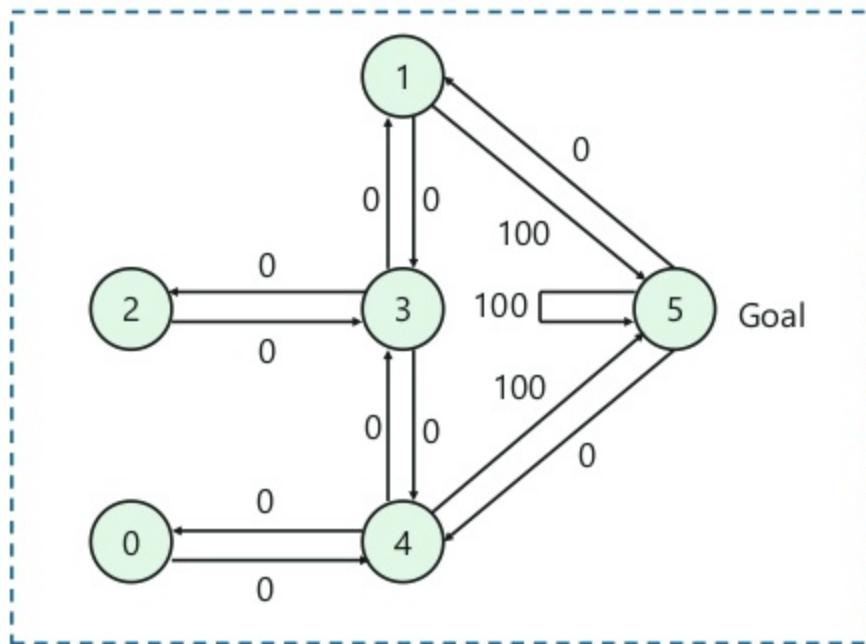


Example (Agent traverse from room 2 to room5):

1. Initial state = state 2
2. State 2 -> state 3
3. State 3 -> state (2, 1, 4)
4. State 4 -> state 5

Understanding Q-Learning With An Example

We can put the state diagram and the instant reward values into a reward table, matrix R .



State	Action					
	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

The -1's in the table represent null values

Understanding Q-Learning With An Example

Add another matrix Q, representing the memory of what the agent has learned through experience.

- The rows of matrix Q represent the current state of the agent
- columns represent the possible actions leading to the next state
- Formula to calculate the Q matrix:

$$Q(state, action) = R(state, action) + \text{Gamma} * \text{Max}[Q(next\ state, all\ actions)]$$

Note

The Gamma parameter has a range of 0 to 1 ($0 \leq \text{Gamma} < 1$).

- If Gamma is closer to zero, the agent will tend to consider only immediate rewards.
- If Gamma is closer to one, the agent will consider future rewards with greater weight

Q – Learning Algorithm

- 1 Set the gamma parameter, and environment rewards in matrix R
- 2 Initialize matrix Q to zero
- 3 Select a random initial state
- 4 Set initial state = current state
- 5 Select one among all possible actions for the current state
- 6 Using this possible action, consider going to the next state
- 7 Get maximum Q value for this next state based on all possible actions
- 8 Compute: $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$
- 9 Repeat above steps until current state = goal state

Q – Learning Example

First step is to set the value of the learning parameter Gamma = 0.8, and the initial state as Room 1.

Next, initialize matrix Q as a zero matrix:

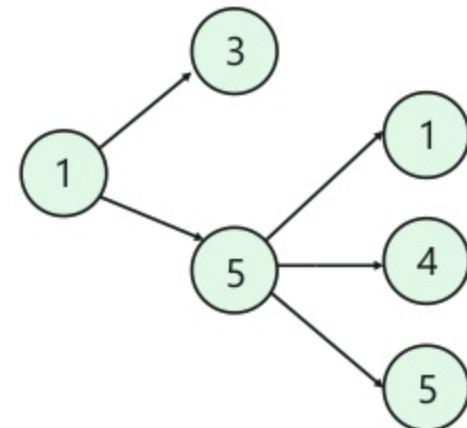
- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum Q value for this next state based on all possible actions:

$$Q(state, action) = R(state, action) + Gamma * \text{Max}[Q(next state, all actions)]$$

$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \left[\begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{matrix} \right] \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \left[\begin{matrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{matrix} \right] \end{matrix}$$



Q – Learning Example

First step is to set the value of the learning parameter Gamma = 0.8, and the initial state as Room 1.

Next, initialize matrix Q as a zero matrix:

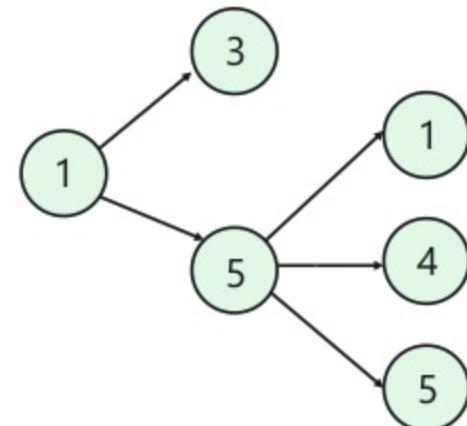
- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum Q value for this next state based on all possible actions

$$Q(state, action) = R(state, action) + Gamma * \text{Max}[Q(next state, all actions)]$$

$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

State	Action					
	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100



Q – Learning Example

For the next episode, we start with a randomly chosen initial state, i.e. state 3

- From room 3 you can either go to room 1,2 or 4, let's select room 1.
- From room 1, calculate maximum Q value for this next state based on all possible actions:

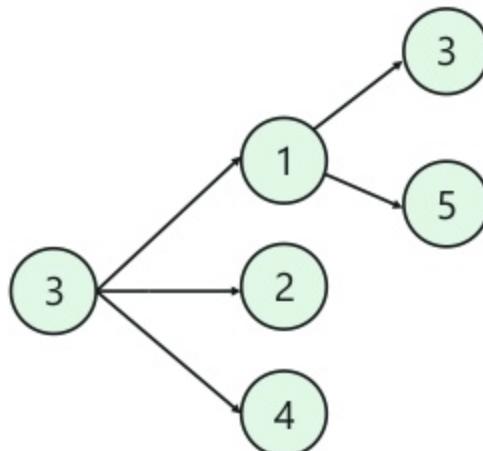
$$Q(state, action) = R(state, action) + \text{Gamma} * \text{Max}[Q(next state, all actions)]$$

$$Q(3,1) = R(3,1) + 0.8 * \text{Max}[Q(1,3), Q(1,5)] = 0 + 0.8 * [0, 100] = 80$$

The matrix Q get's updated

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	80	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

State	Action					
	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100



Q – Learning Example

For the next episode, the next state, 1, now becomes the current state. We repeat the inner loop of the Q learning algorithm because state 1 is not the goal state.

- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum Q value for this next state based on all possible actions:

$$Q(state, action) = R(state, action) + \text{Gamma} * \text{Max}[Q(next state, all actions)]$$

$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

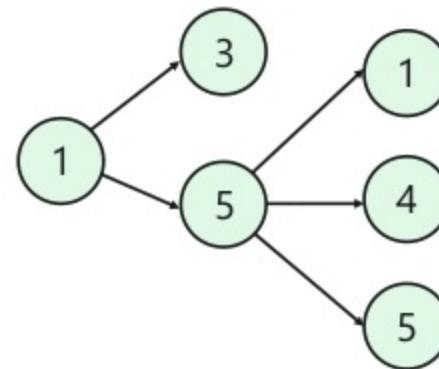
The matrix Q remains the same since, Q(1,5) is already fed to the agent

State	Action					
	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	100
2	0	0	0	0	0	0
3	0	80	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

$Q =$

$$R =$$

State	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100



Demo



Thank You

For more information please visit our website
www.edureka.co